

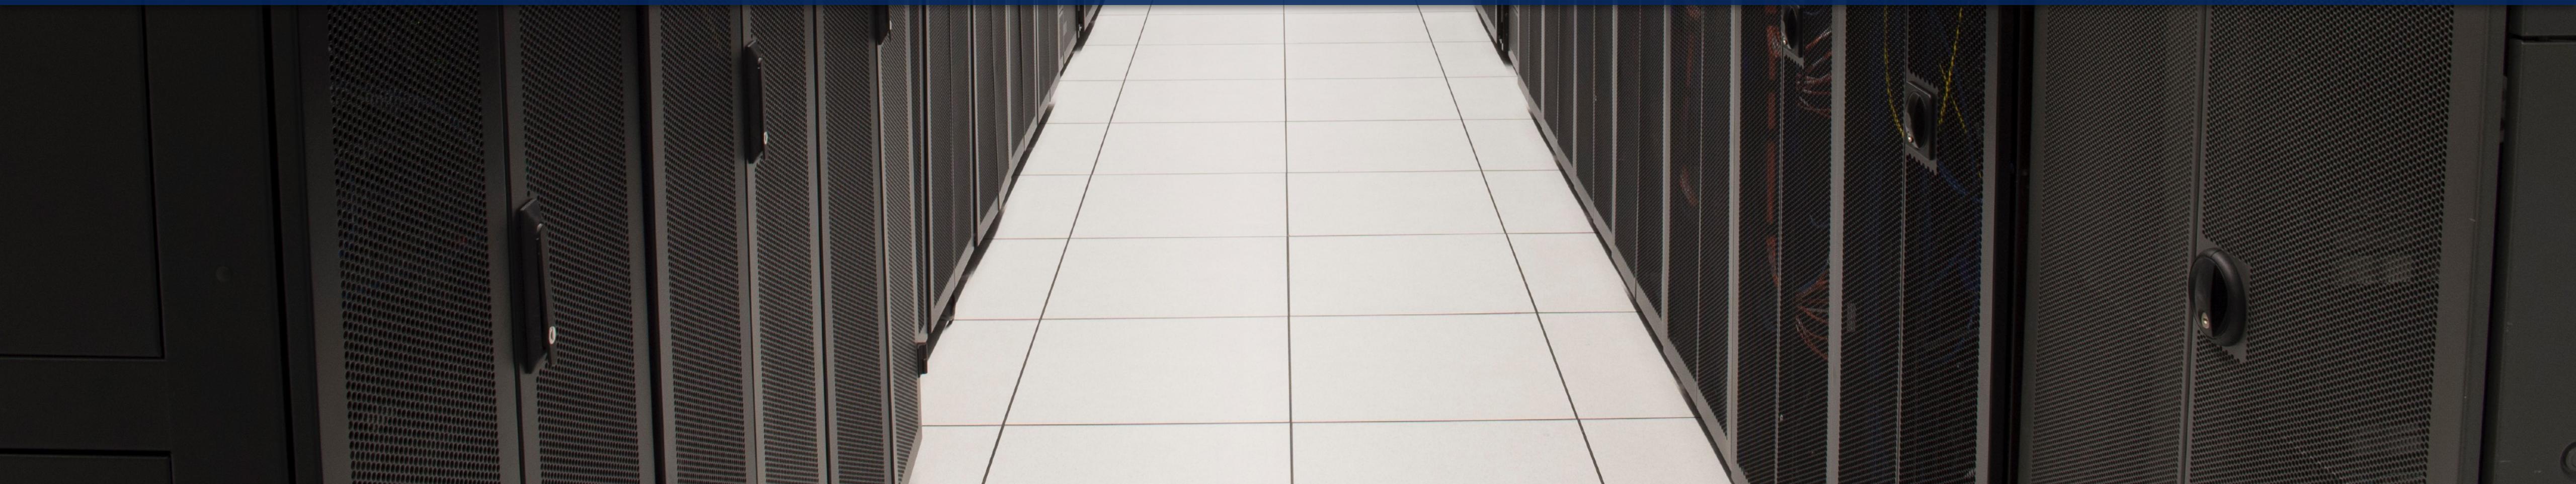
ManeFrame II (M2) Introduction

Dr. Robert Kalesky
HPC Applications Scientist
Office of Information Technology and Center for Research Computing

Dr. John LaGrone
HPC Applications Scientist
Office of Information Technology and Center for Research Computing



OIT Research and Data Science Services Center for Research Computing



Research and Data Science Services



- Provide research computing support
- High-Performance Computing - Dr. Robert Kalescky
- High-Performance Computing - Dr. John LaGrone
- Data Science - Dr. Eric Godat
- Data Science - Dr. Tue Vue
- Custom Devices (IOT, wearables, etc.) - Guillermo Vasquez

Center for Research Computing



- Maintains our primary shared resource for research computing, ManeFrame II (M2), in collaboration with OIT
- Provides research computing tools, support, and training to all faculty, staff, and students using research computing resources
- www.smu.edu/crc has documentation and news
- help@smu.edu or rkalescky@smu.edu for help
- Request an account at www.smu.edu/crc

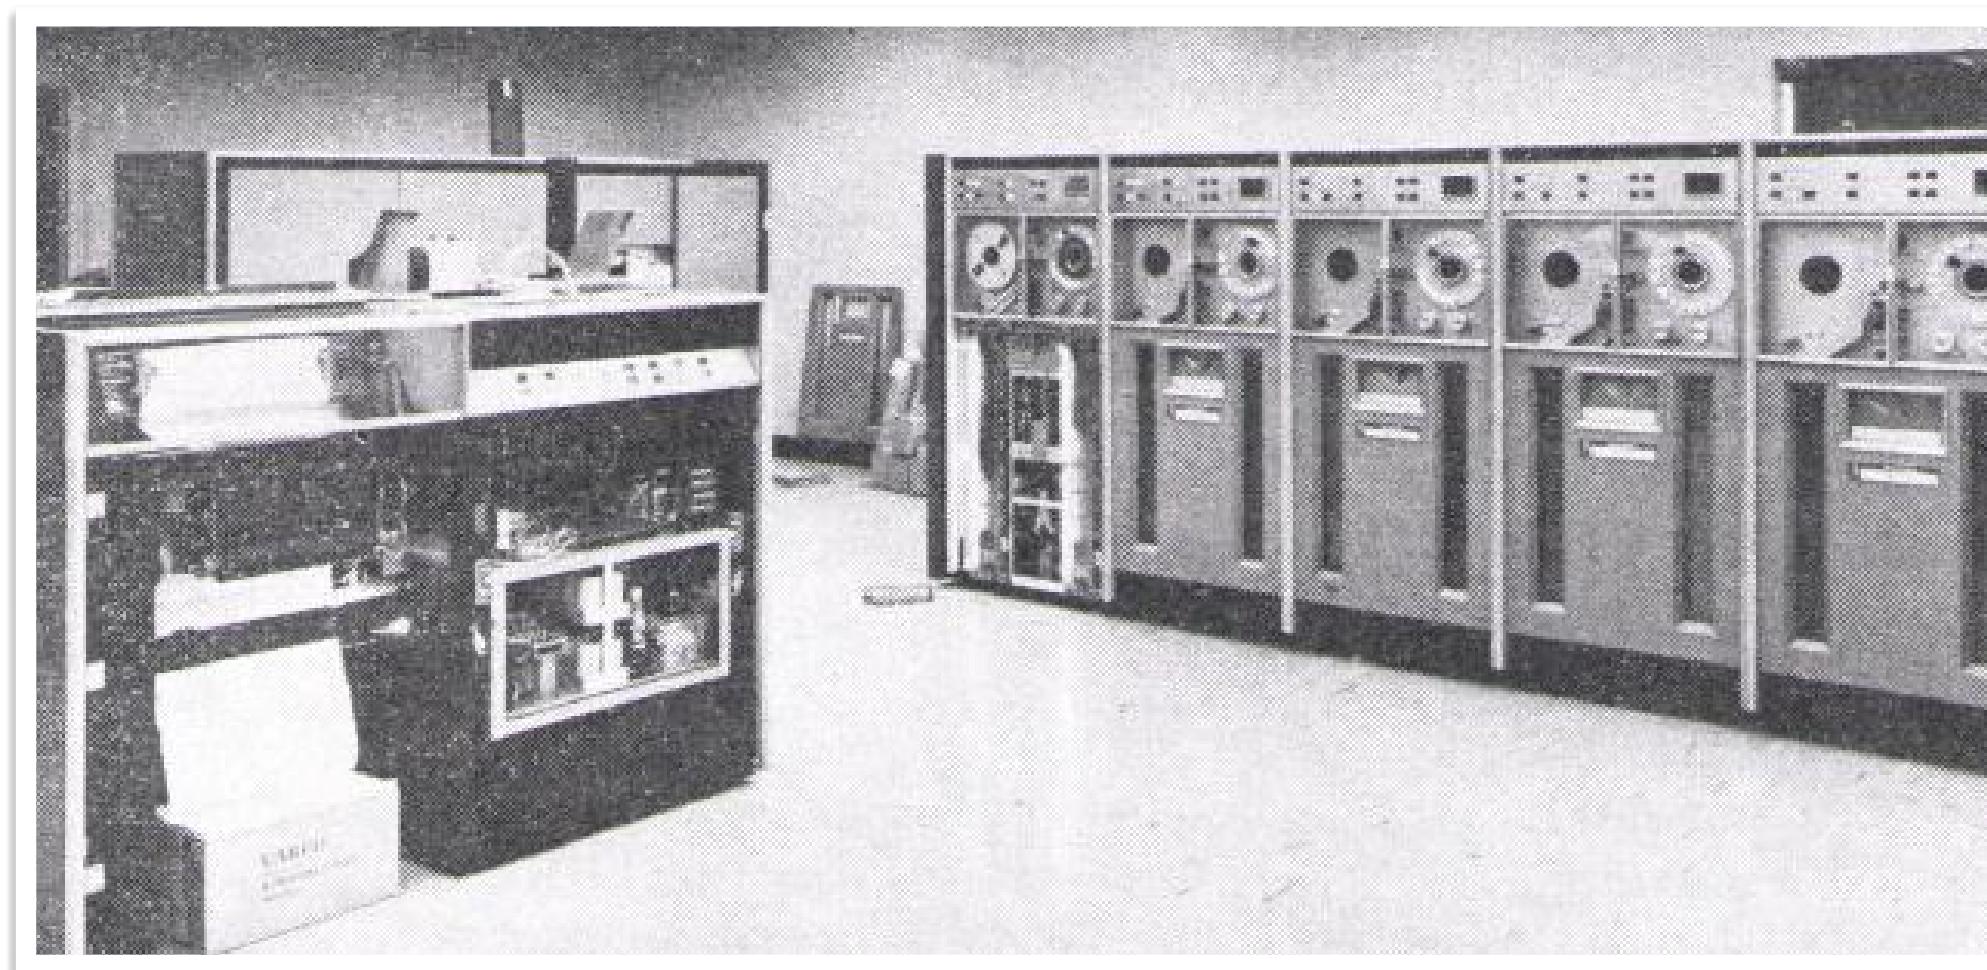
SMU's HPC History



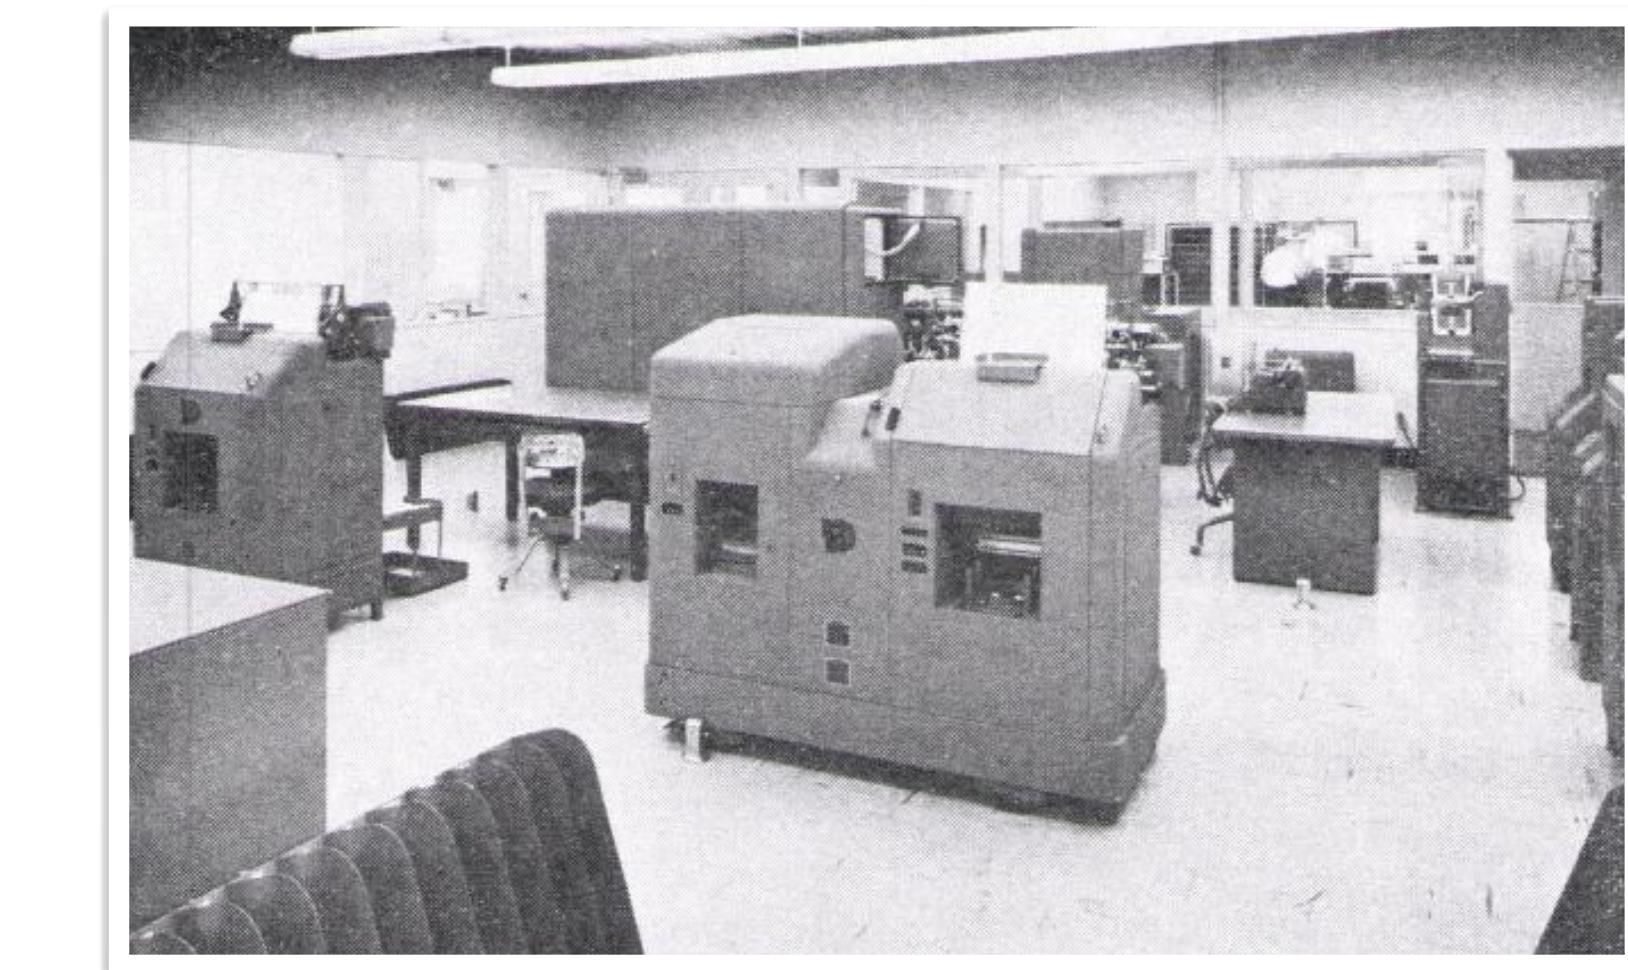
UNIVAC Computing Center in 1957



Punch card machines



Printer (left) and magnetic tape storage (right)



UNIVAC 120 (back left)

High Performance Computing





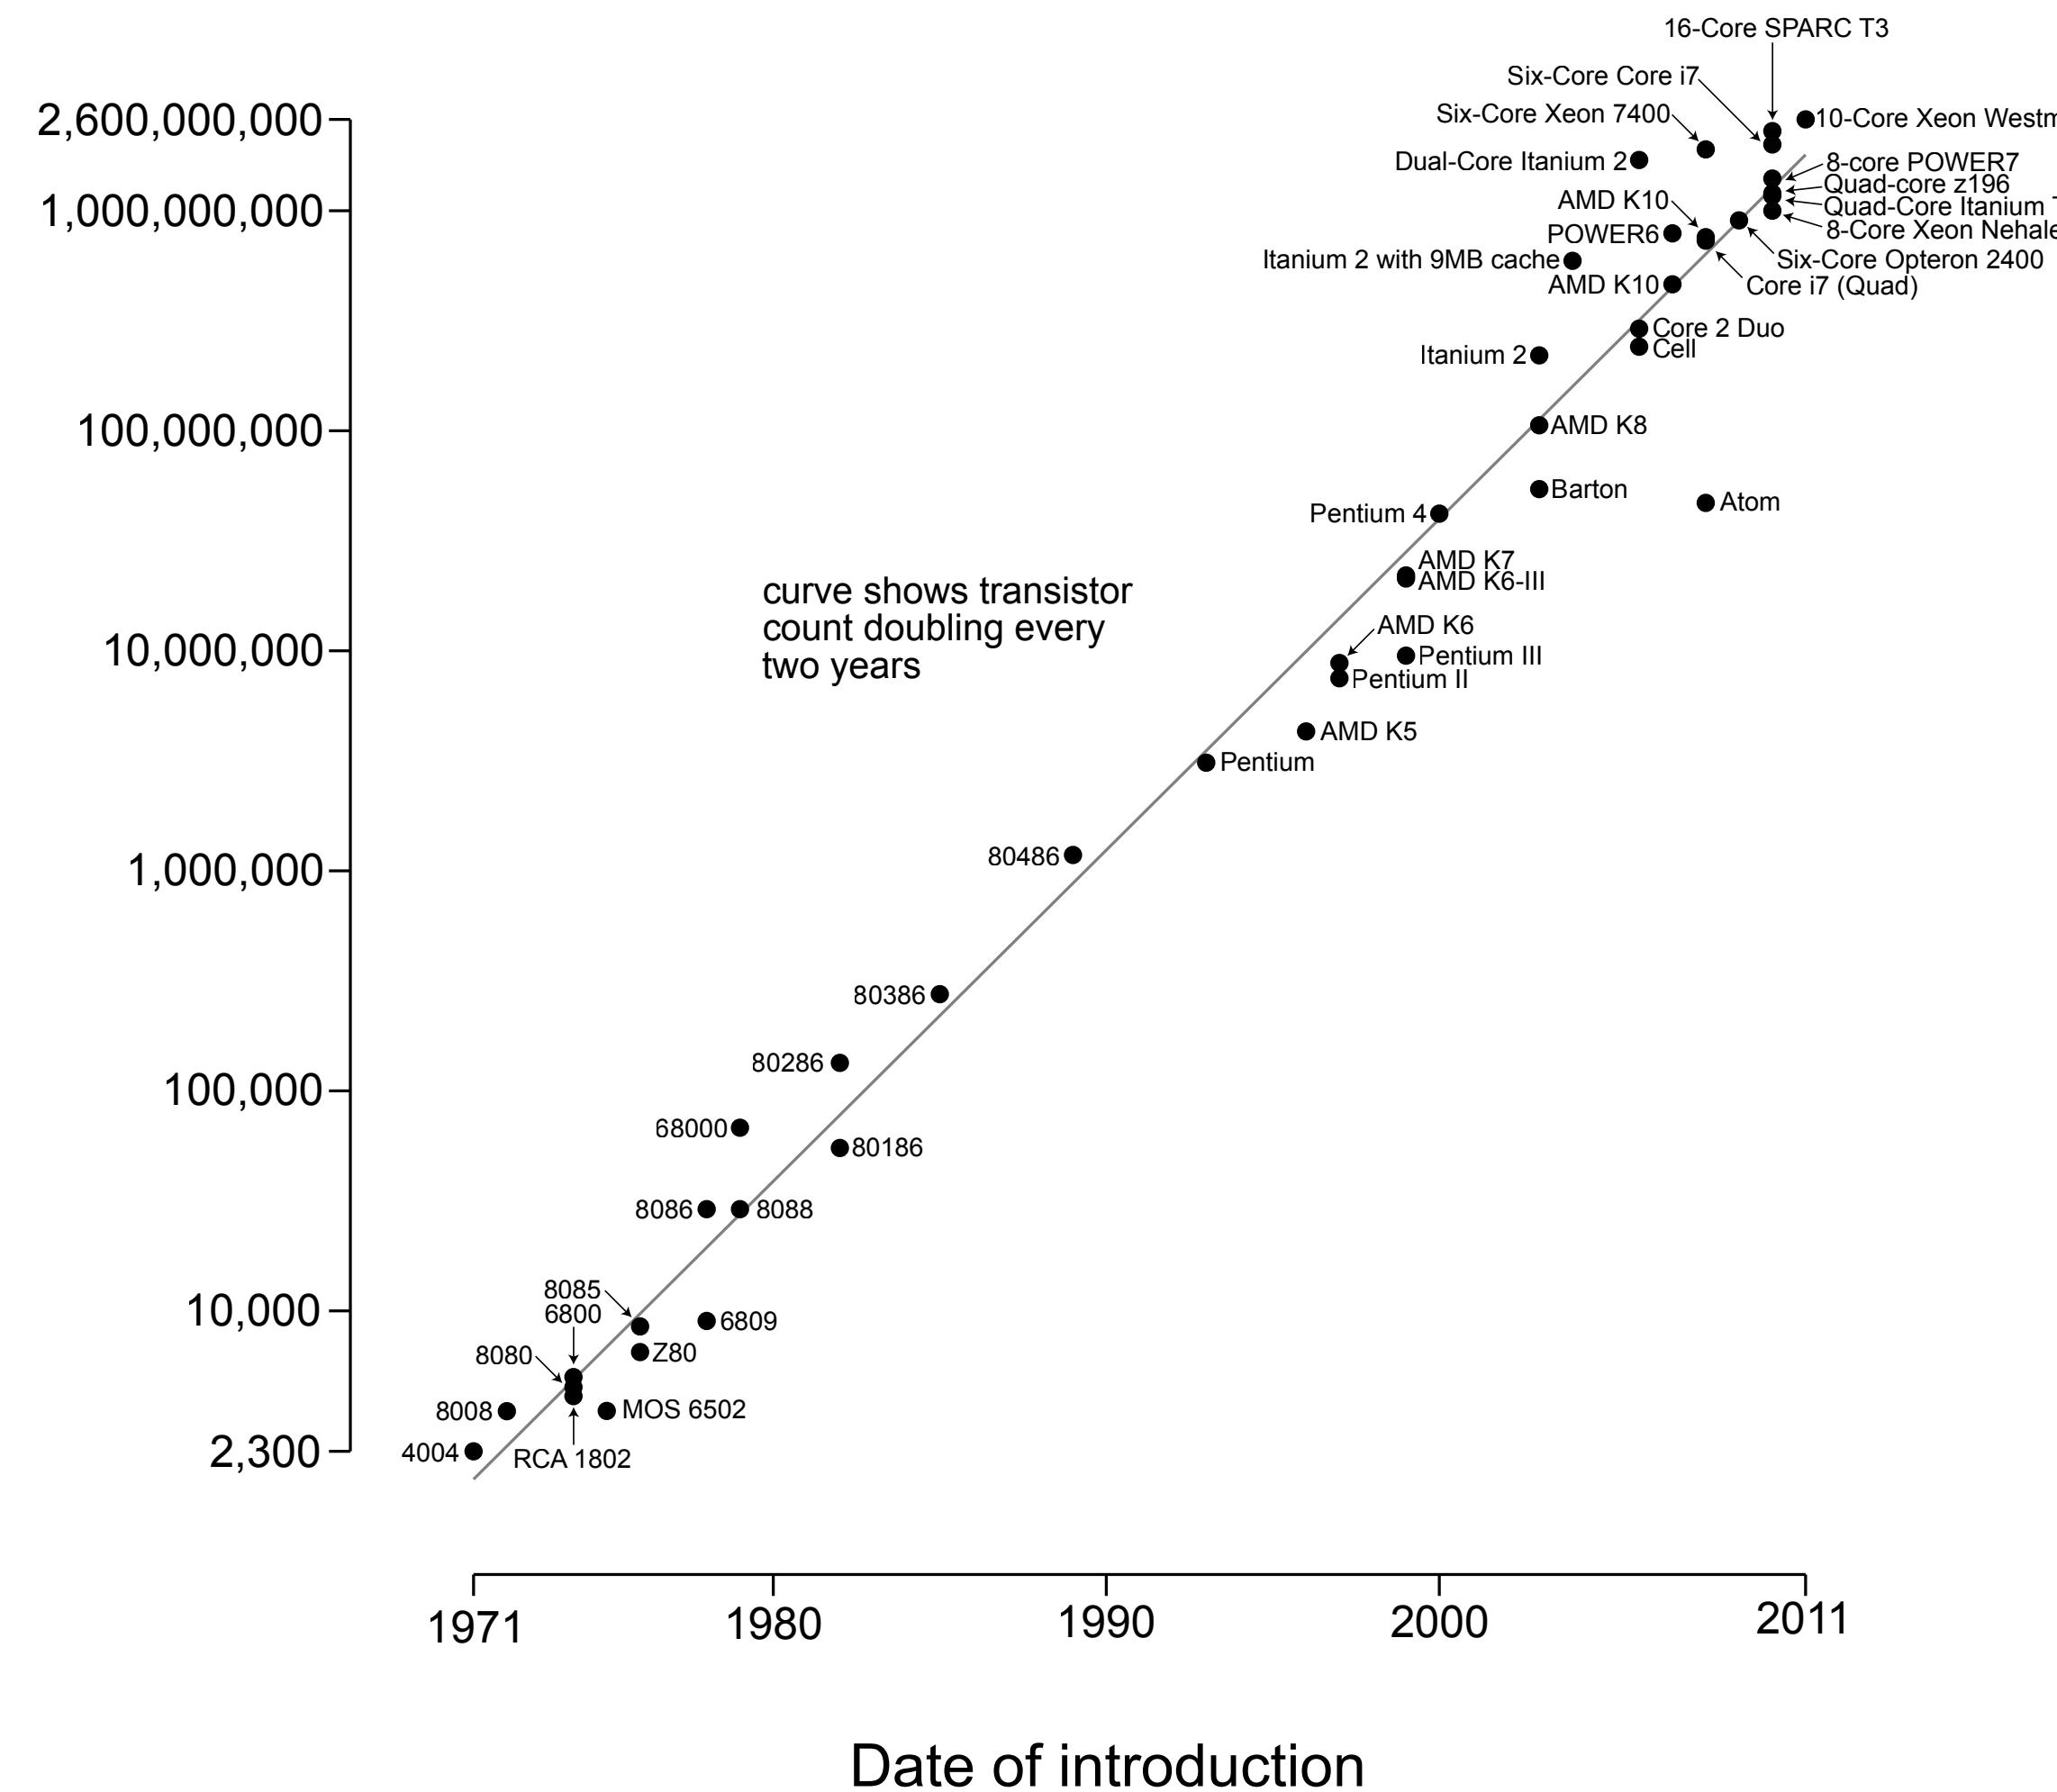
Moore's Law

- Historically, we have depended on hardware advances to enable faster and larger simulations
- In 1965, Gordon Moore observed that the CPU and RAM transistor count about doubled each year
- “Moore’s Law” has since been revised to a doubling once every 2 years, with startling accuracy
- Physical limits, e.g. power consumption, heat emission, and even the size of the atom, have currently stopped this expansion on individual processors, with speeds that have leveled off since around 2008



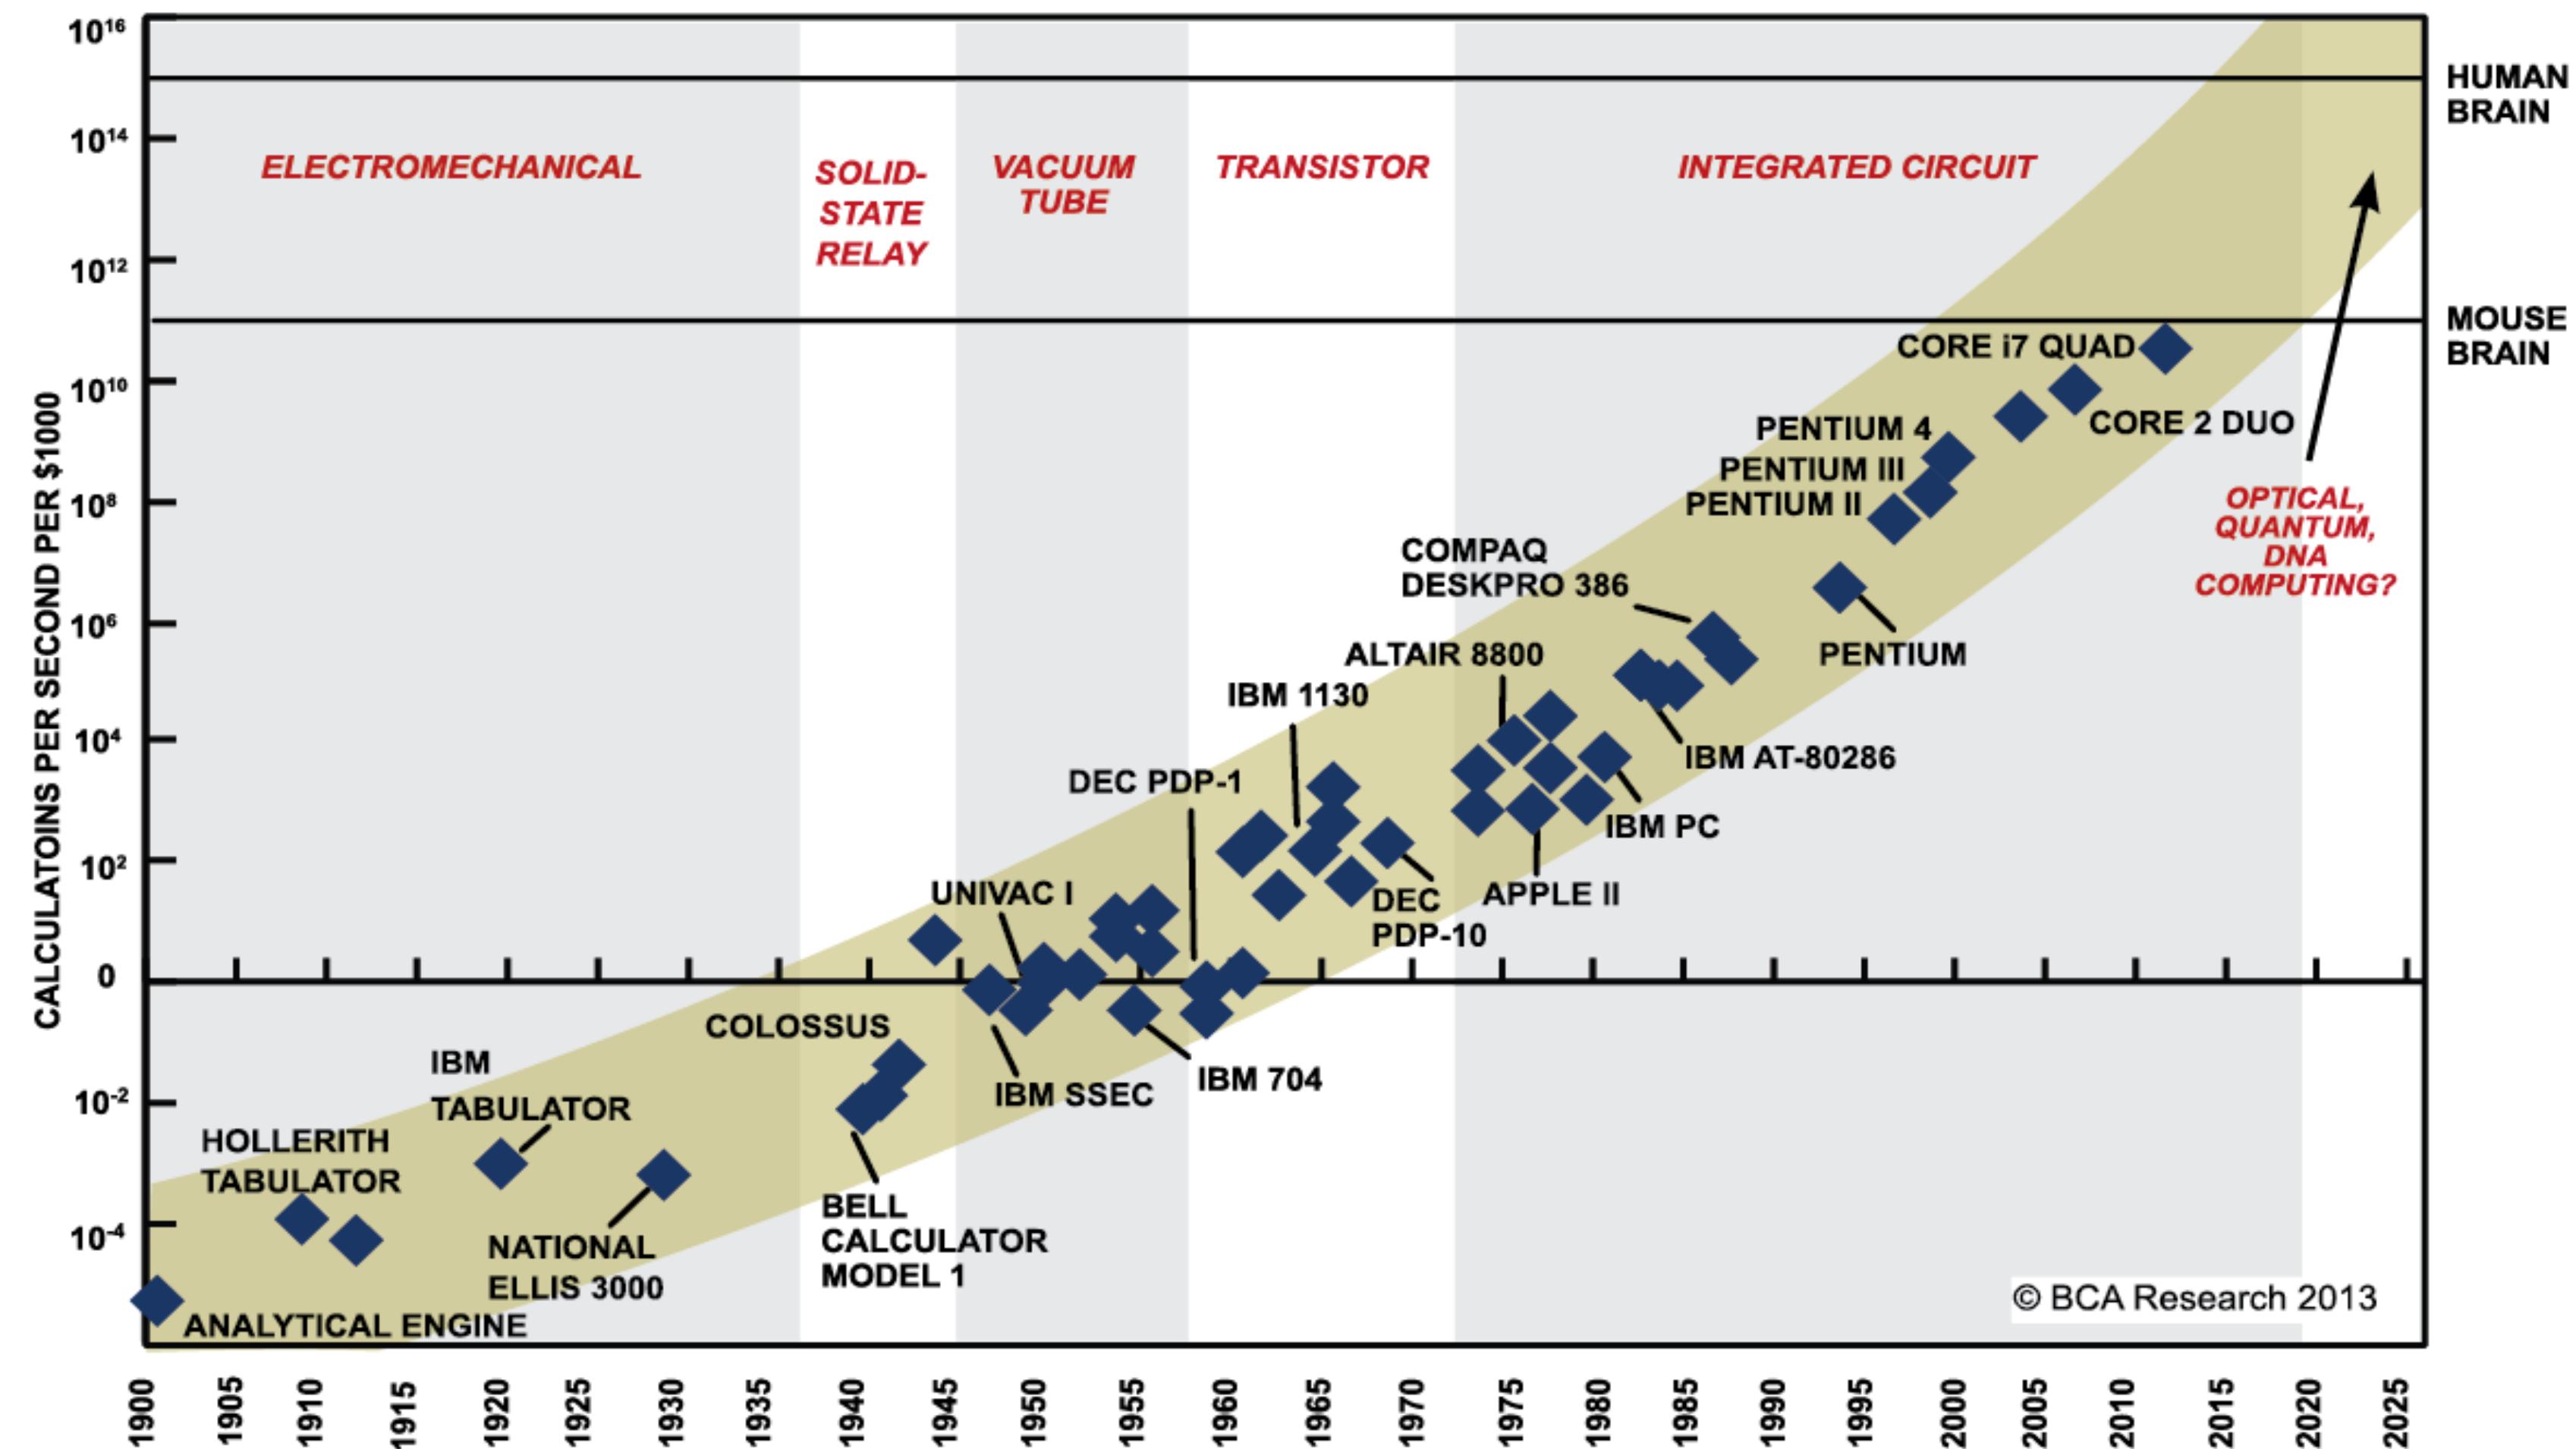
Growth in CPU Transistor Counts

Microprocessor Transistor Counts 1971-2011 & Moore's Law





Decrease in Cost Per FLOPS



SOURCE: RAY KURZWEIL, "THE SINGULARITY IS NEAR: WHEN HUMANS TRANSCEND BIOLOGY", P.67, THE VIKING PRESS, 2006. DATAPoints BETWEEN 2000 AND 2012 REPRESENT BCA ESTIMATES.



Additional Problems

- Growing performance discrepancy between processing units and data storage
 - 40% processor performance improvement per year
 - 10% RAM performance improvement per year
 - Discrepancy leads to memory-bound programs
 - CPU spends more and more time idle, waiting on data from
- Many simulations require incredible amounts of memory to achieve high-accuracy solutions (PDE, MD, QM, etc.)
 - These cannot fit on a single computer alone

The Parallel Solution



- Using multiple processing units (cores, etc.) simultaneously to perform a computation
- Use multiple computers to store data for large problems
- Essentially all modern CPUs have multiple cores



Unfortunately...

“I know how to make 4 horses pull a cart - I don't know how to make 1024 chickens do it.”

Enrico Clementi

What is High Performance Computing?



High performance computing (HPC) is the use of computing resources that are significantly more powerful than what is commonly available.*

*As such, it's always a moving target.

Supercomputer Types

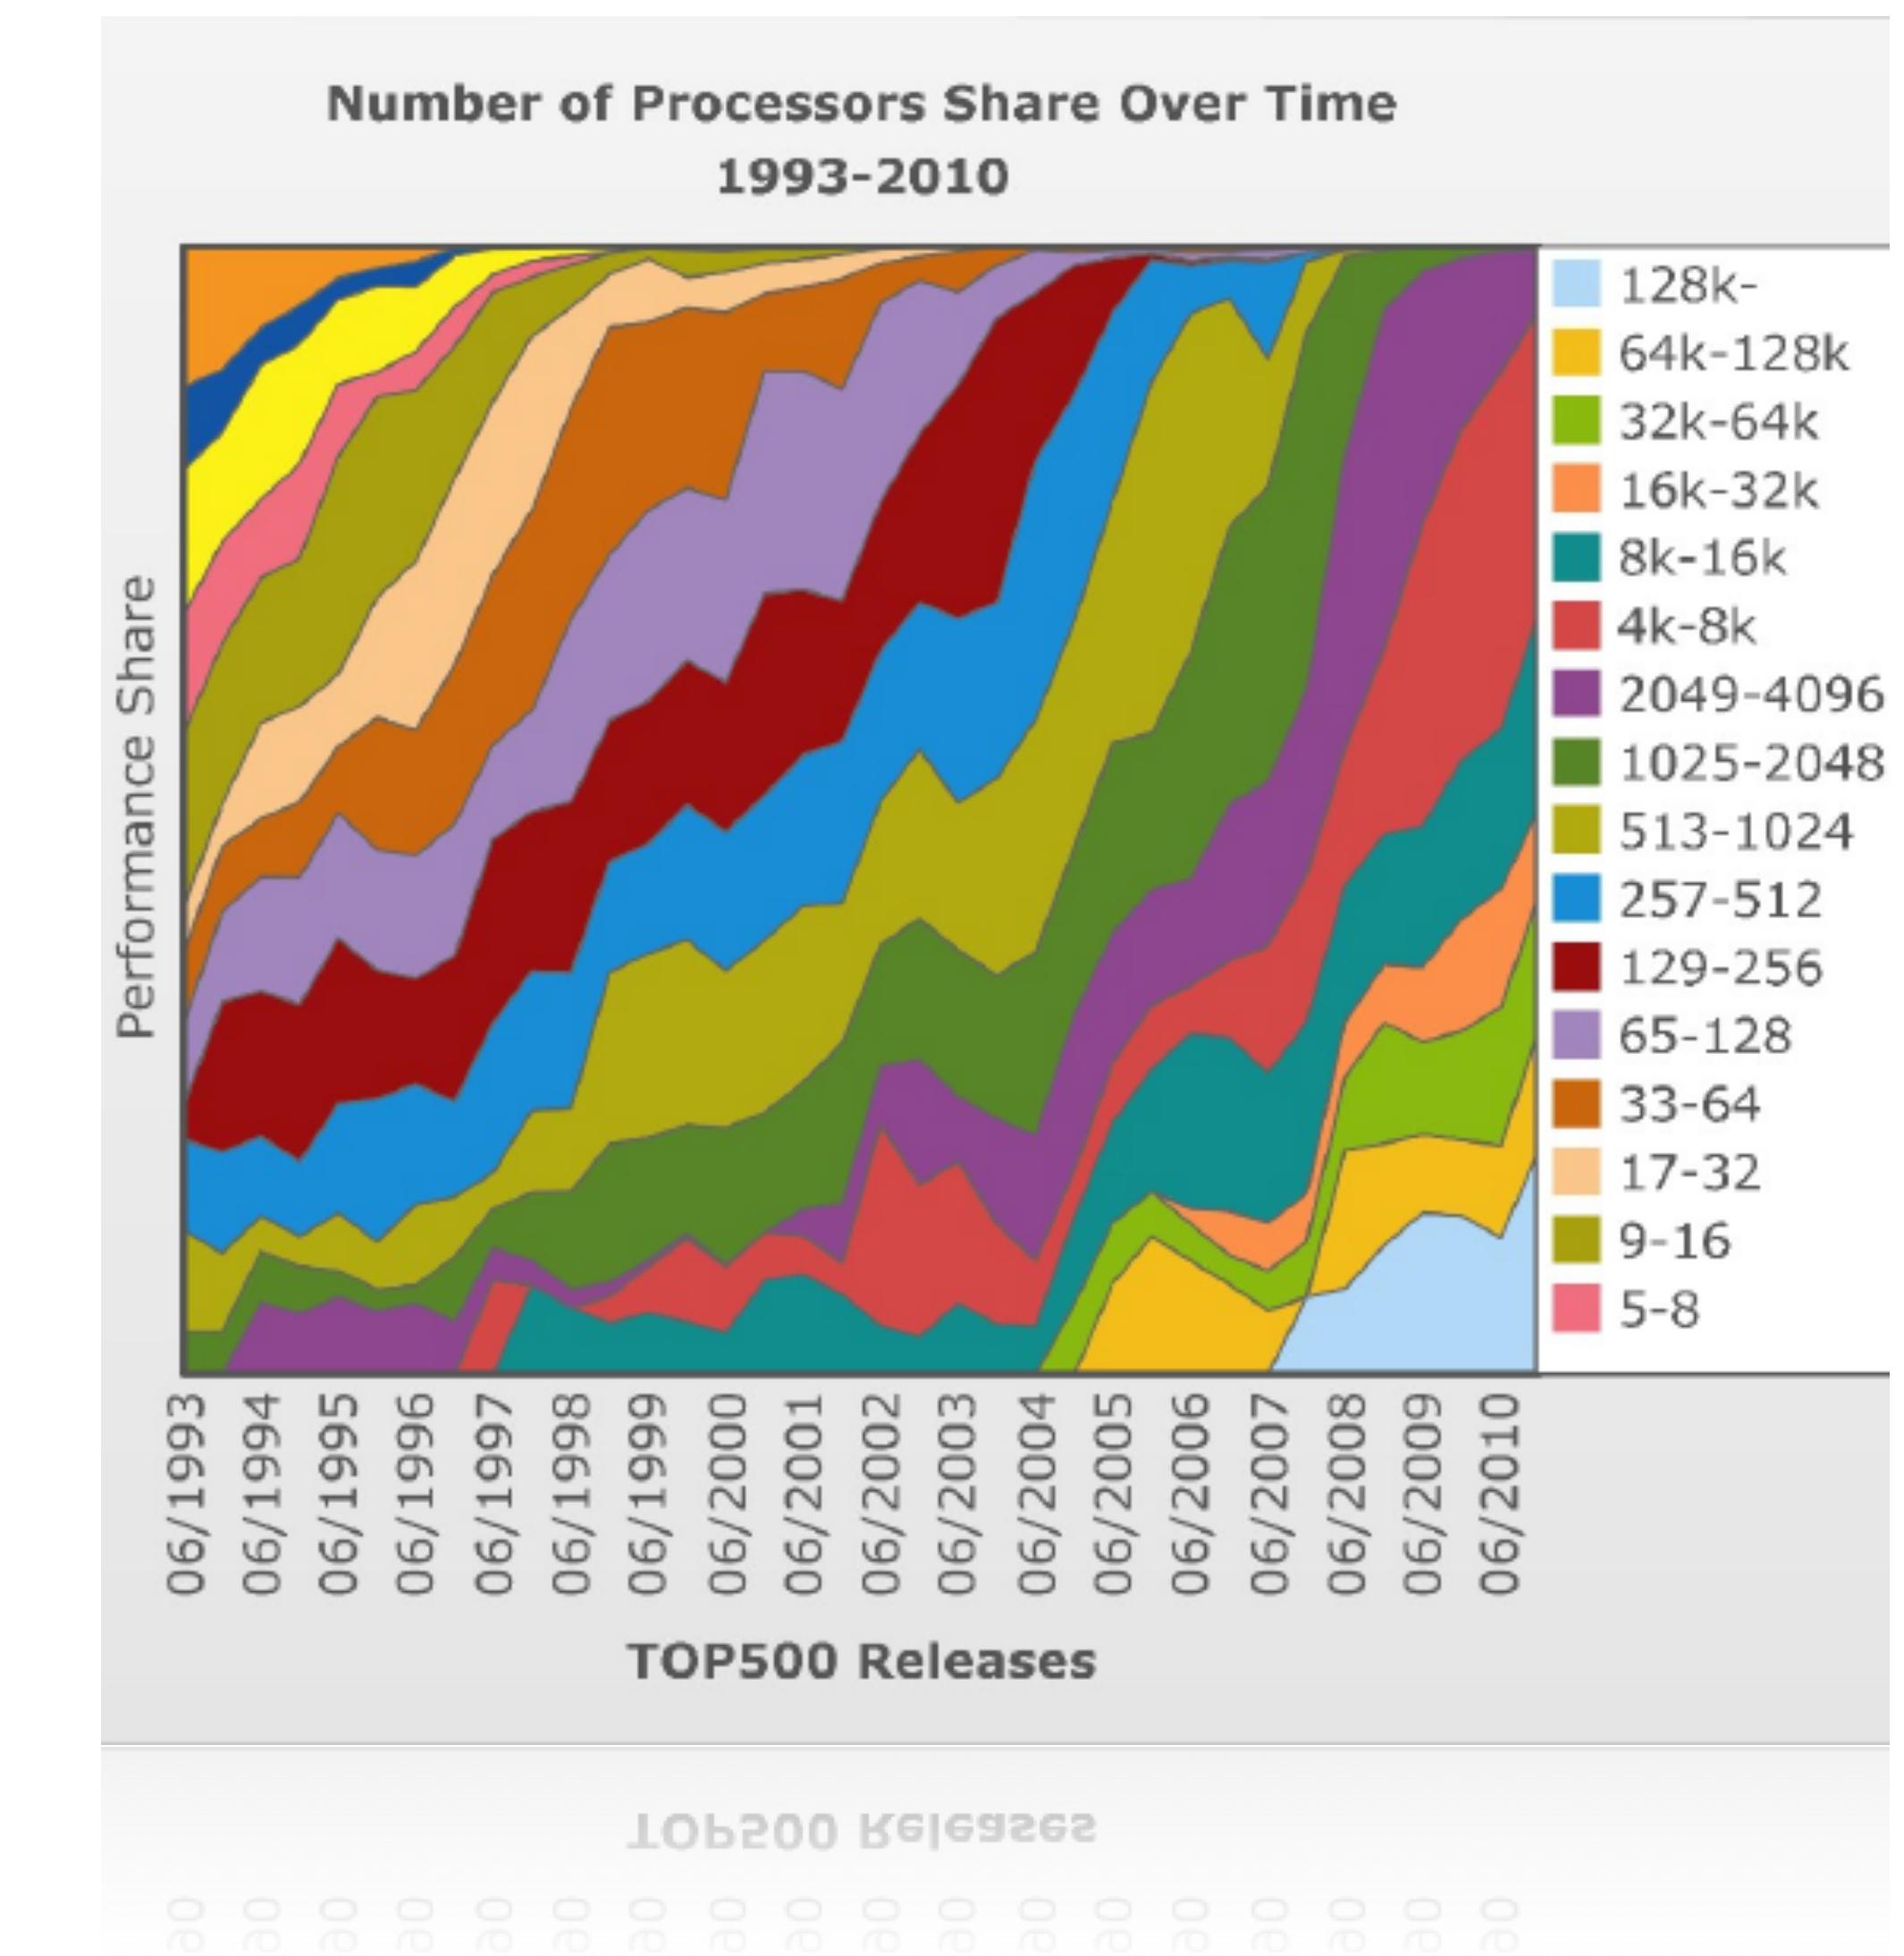
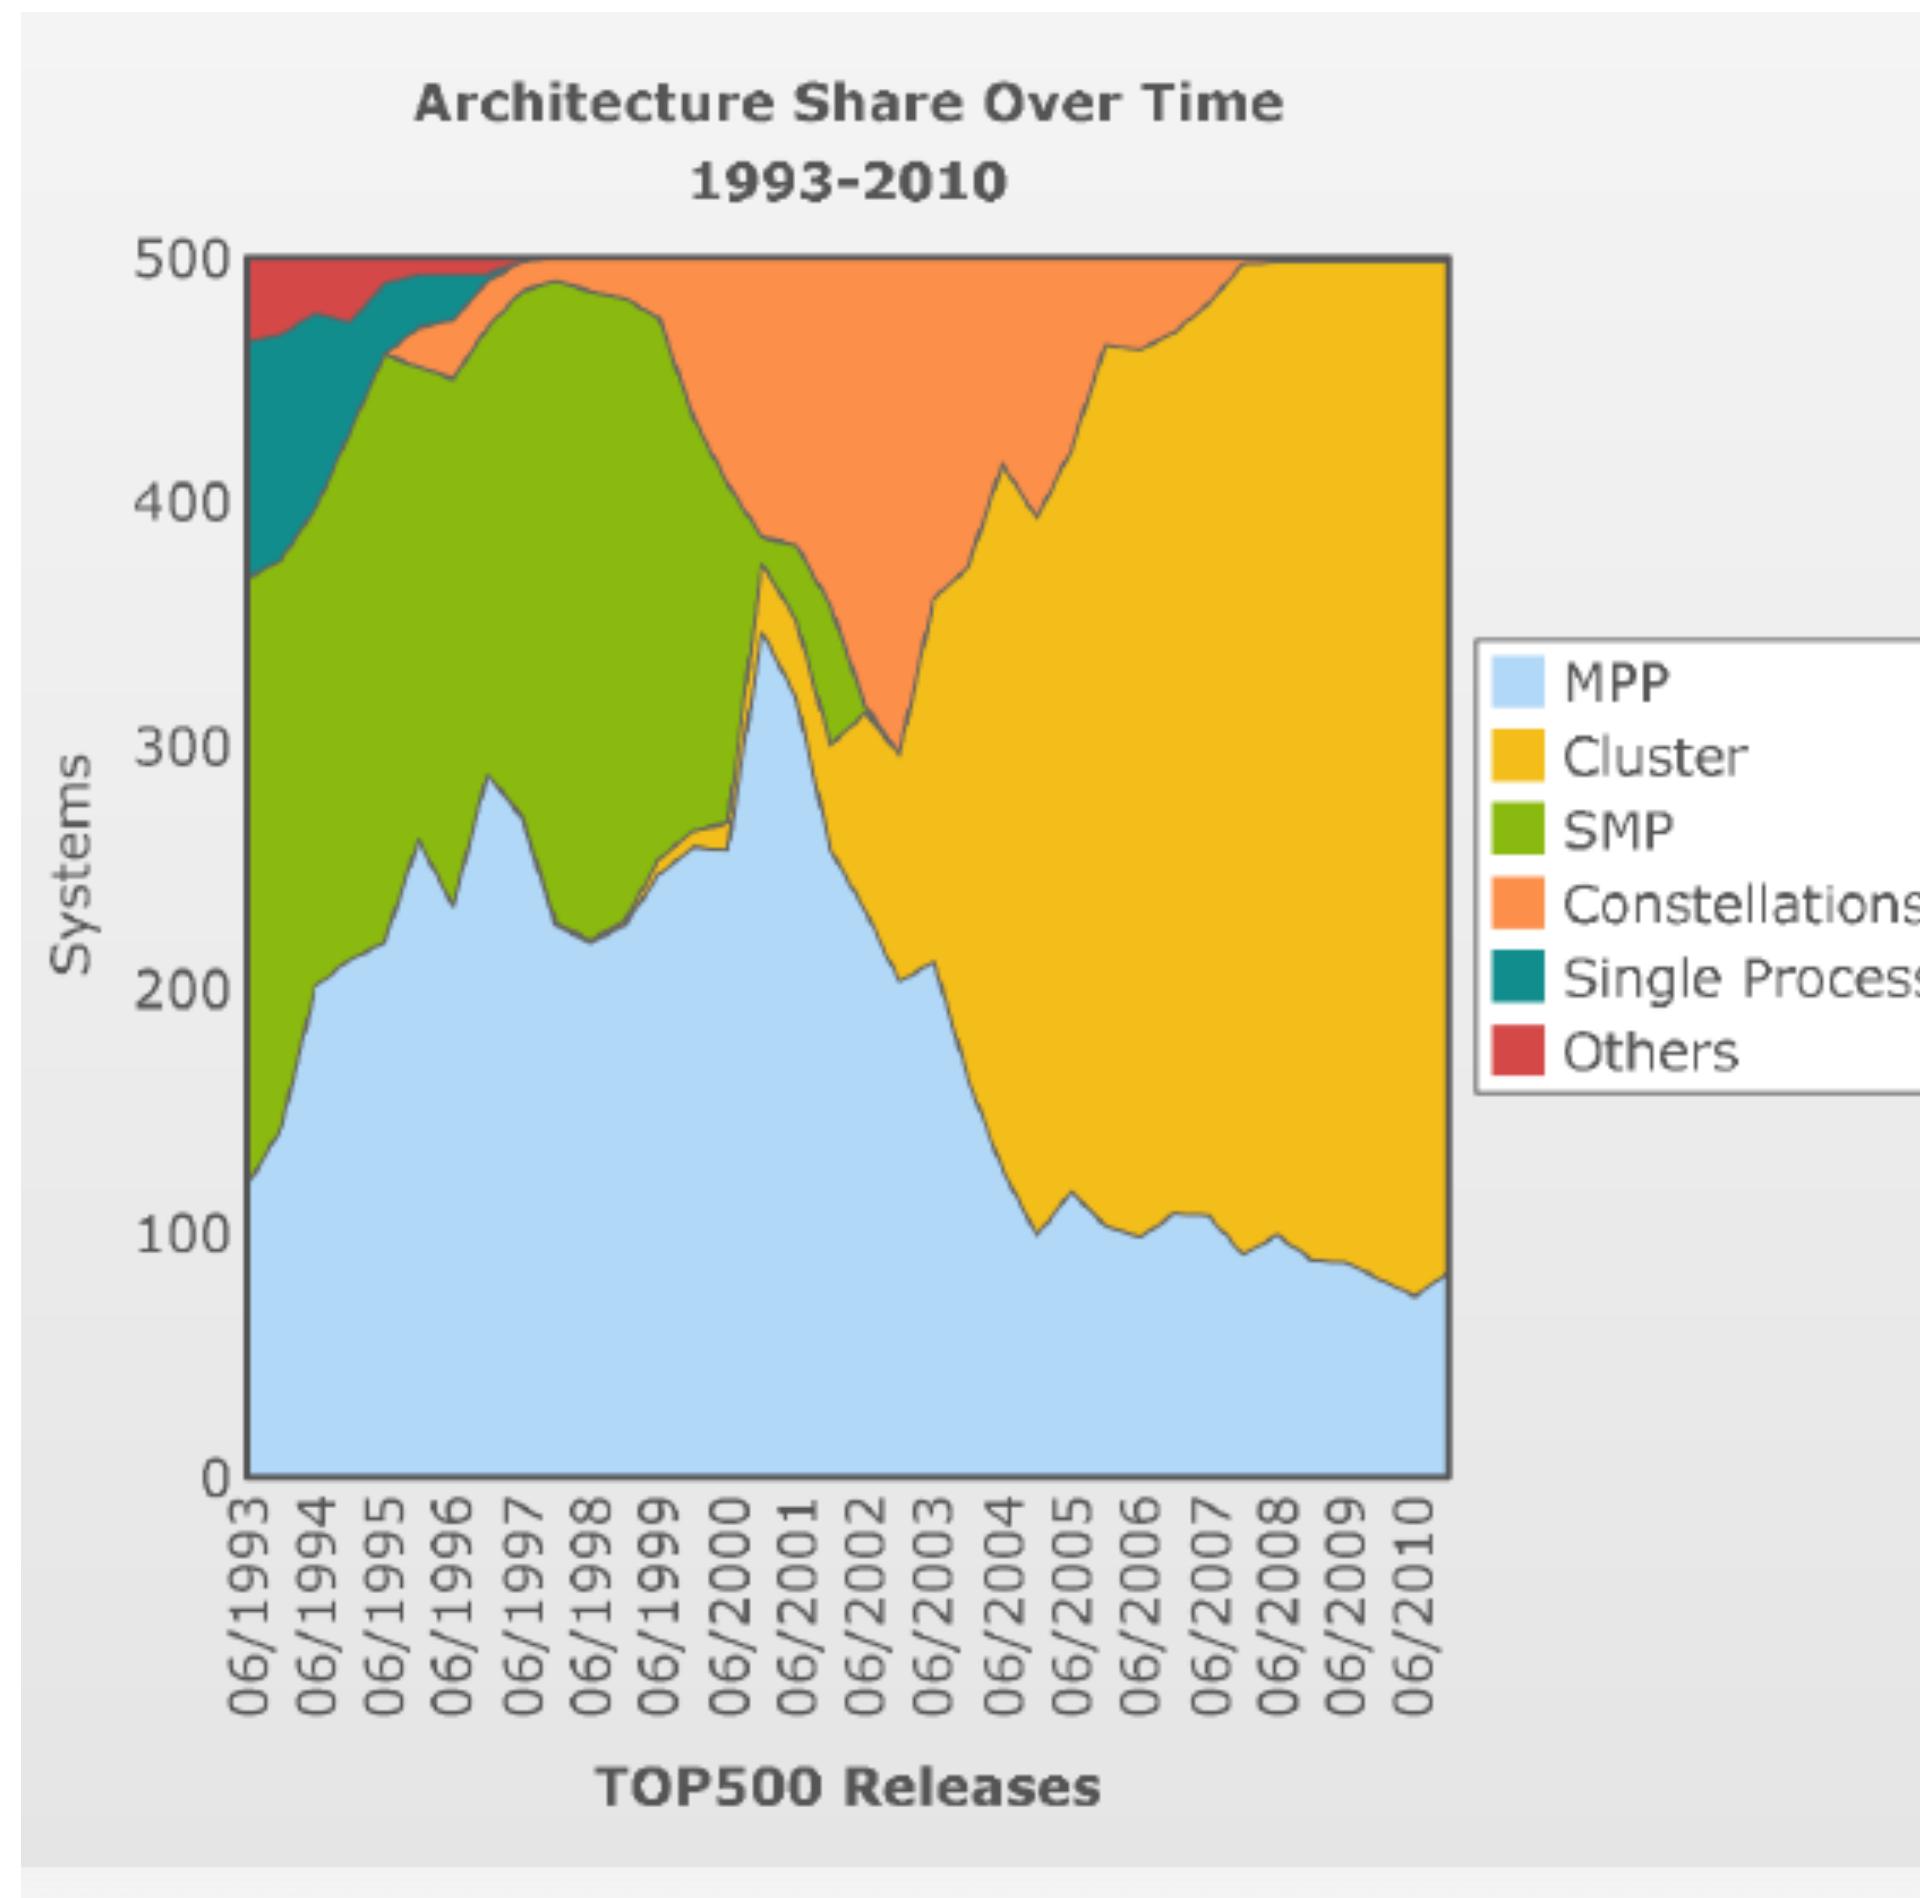


- Single processor
- Shared Memory Parallel (SMP)
- Massively Parallel Processors
- Constellations
- Clusters

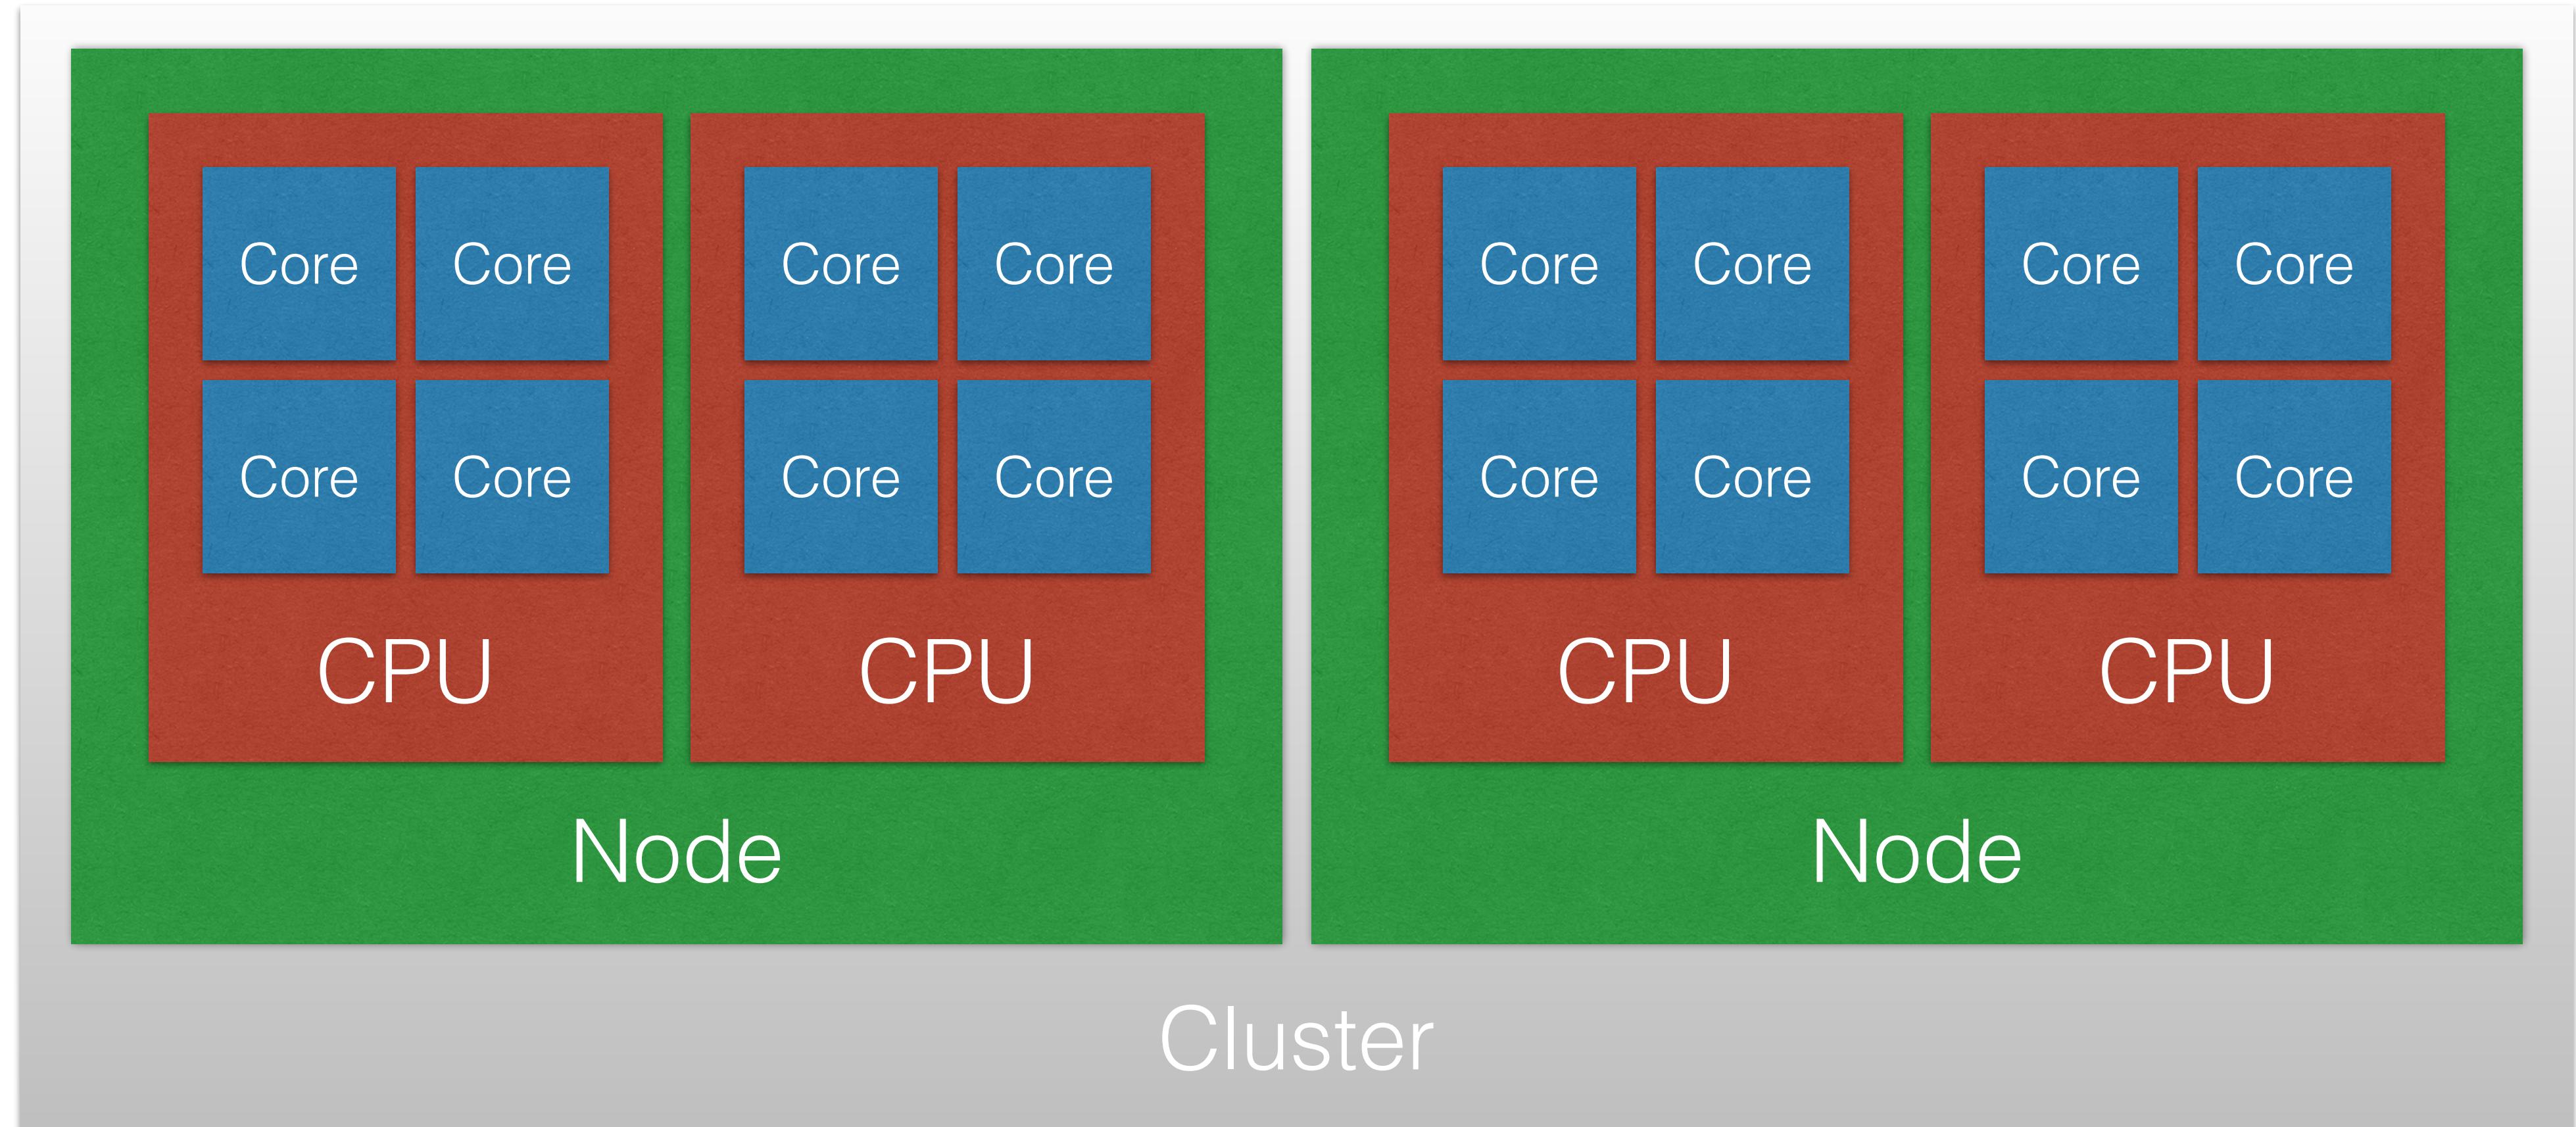




History of Parallel Architectures



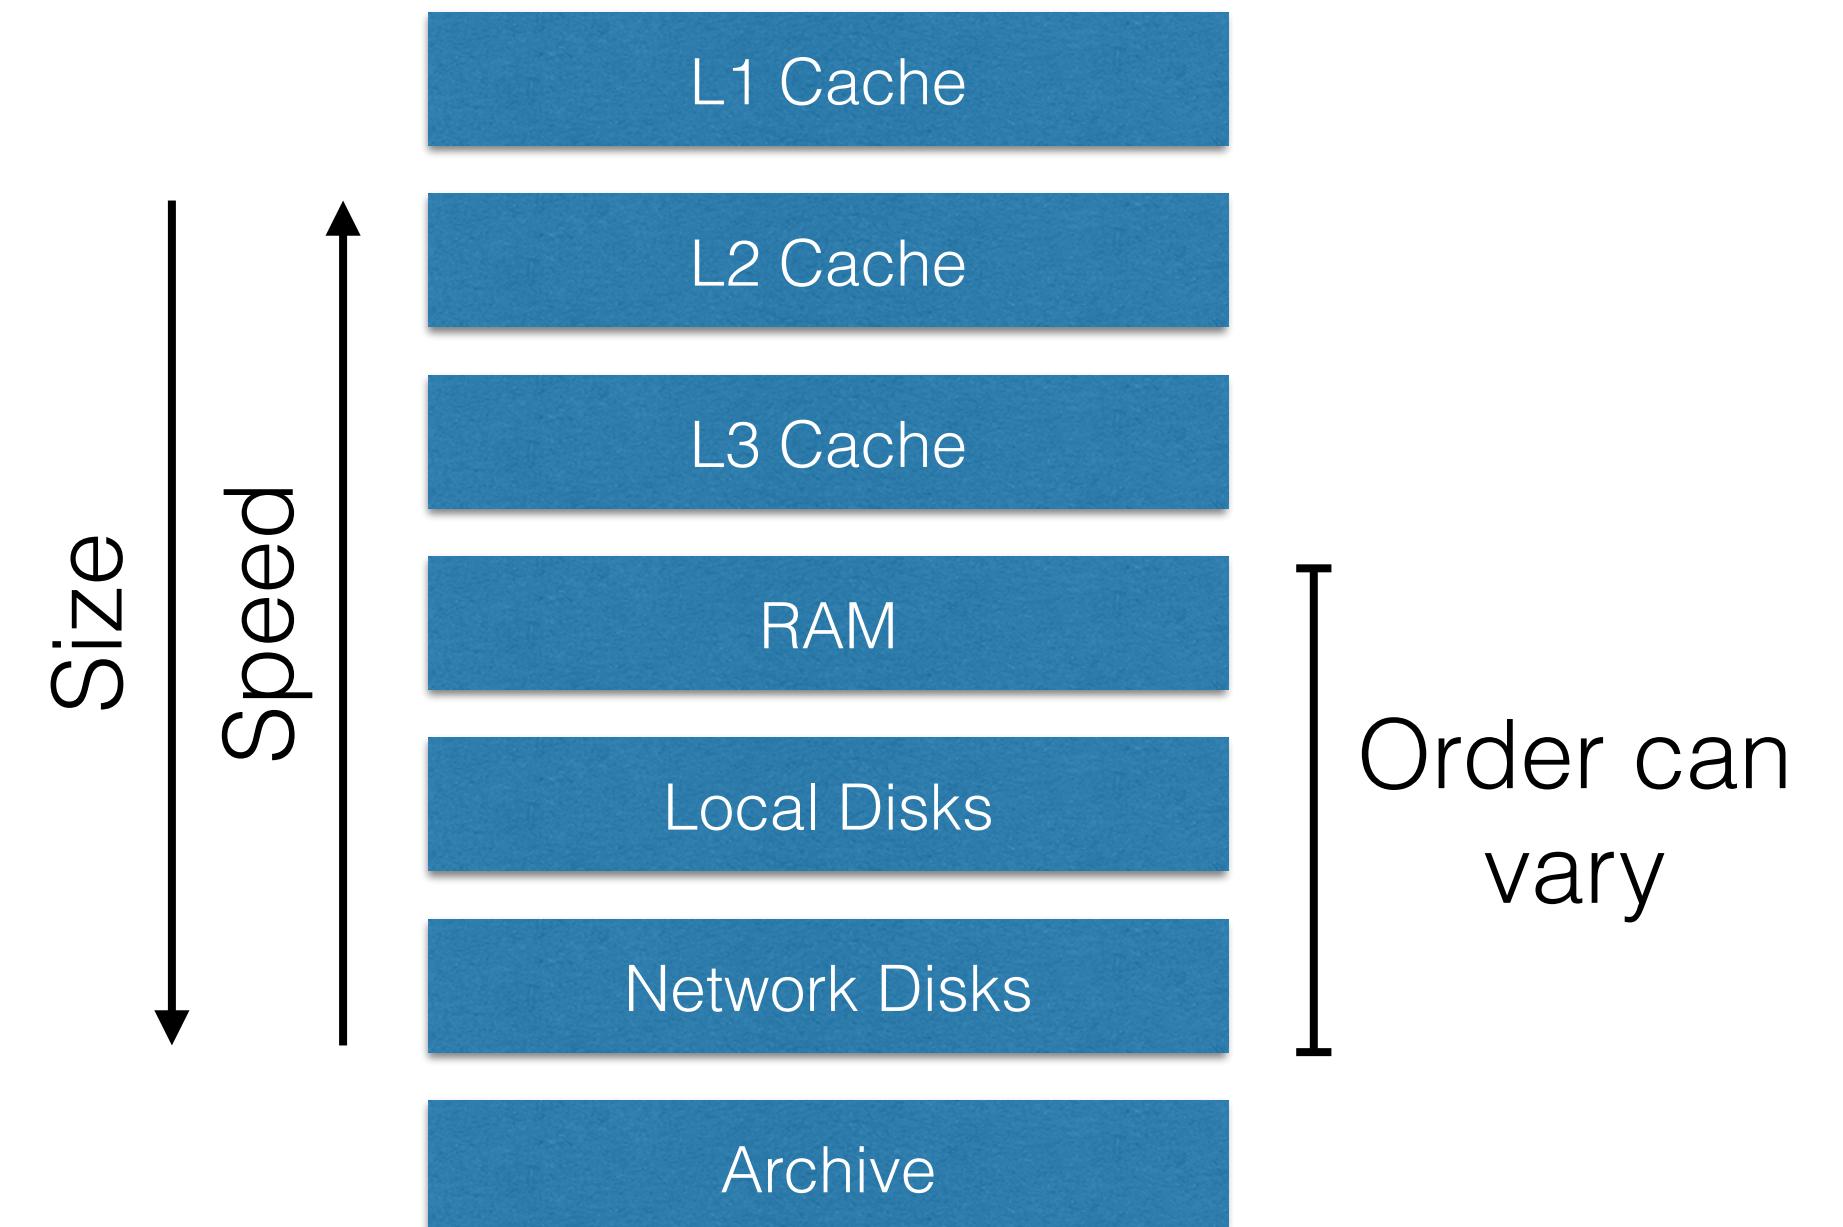
Cluster Compute Resources

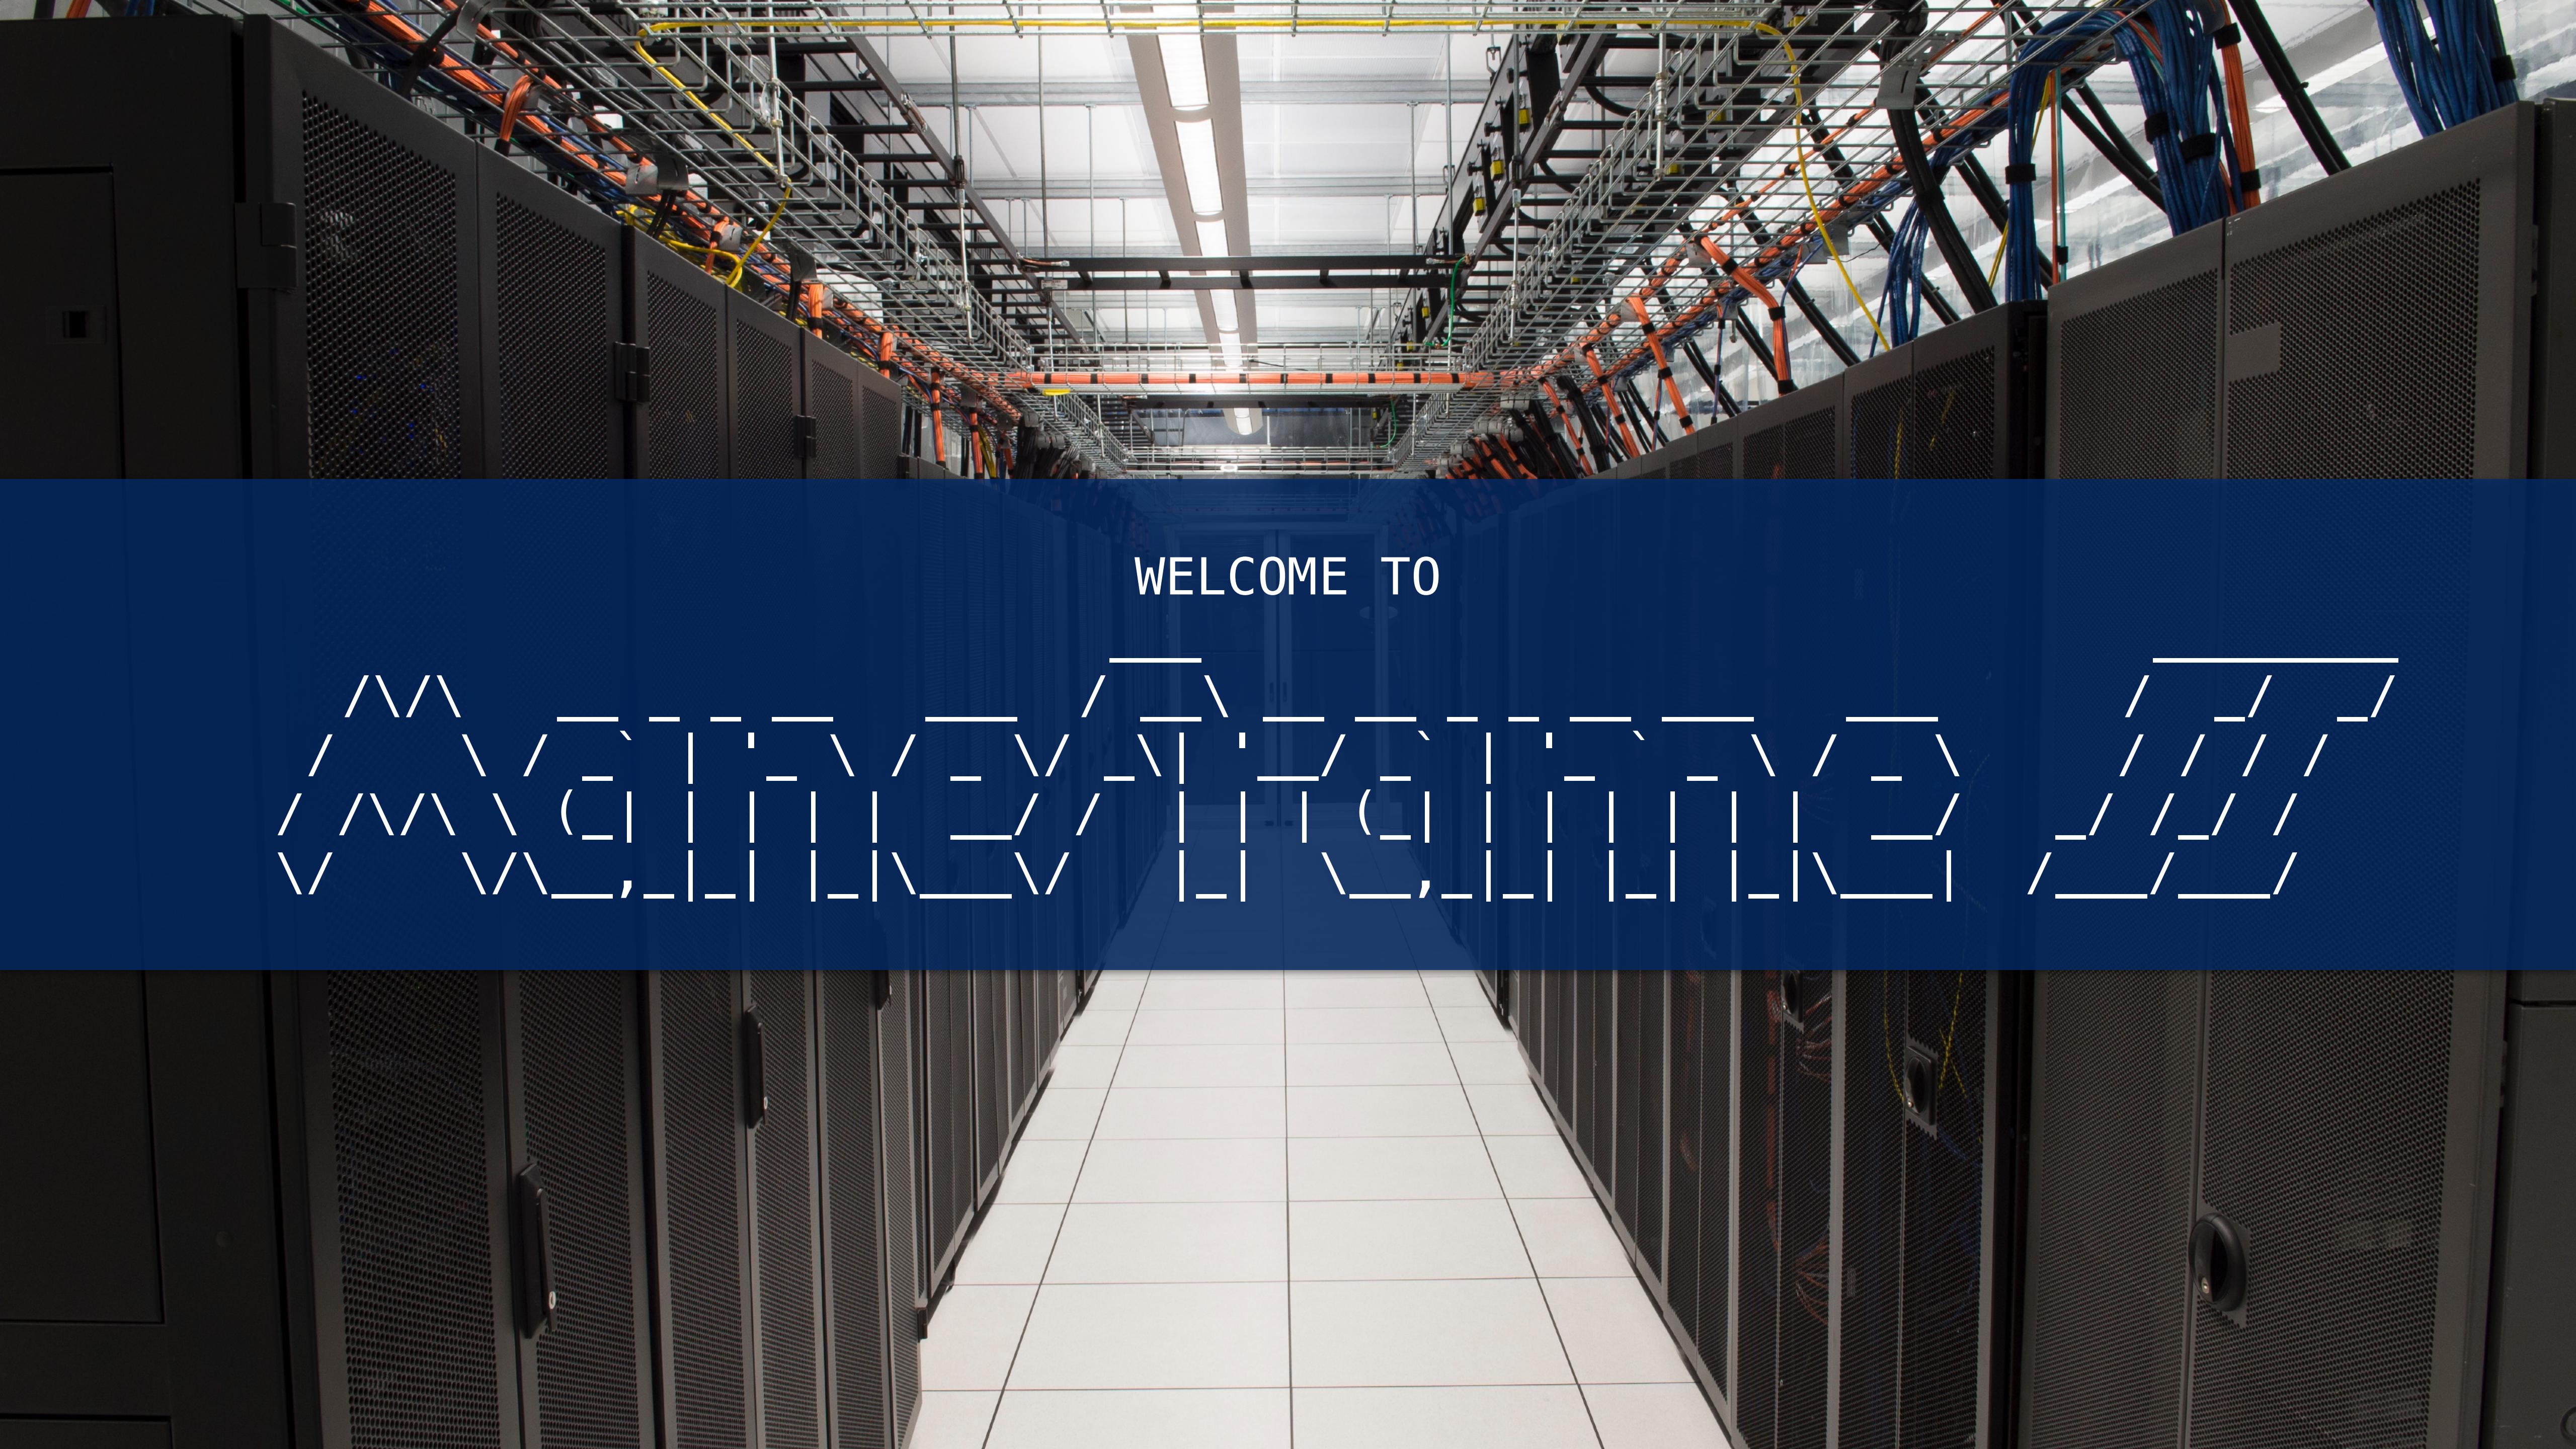




Cluster Storage Resources

- Usually there is an inverse relationship between storage amount and speed
- Clusters and applications can distort the standard order
 - Aggregation, e.g. RAID
 - Network optimization to a specific topology
 - Applications can be configured and built for specific hardware
 - New hardware technologies





WELCOME TO



ManeFrame II



- High performance features:
 - CPUs
 - Accelerators
 - Networking technologies
 - Substantial memory per node
 - Interactive GPU-accelerated remote desktop experiences



ManeFrame II

- CentOS 7 operating system
- SLURM resource scheduler
- Lmod environment module system.
- Development tool chains including the GCC, LLVM, Intel, and PGI compiler suites
- Optimized high-level programming environments such as MATLAB, Python, and R

ManeFrame II Node Types



Type	Quantity	Cores	Memory [GB]	Additional Resources
Standard-Memory	176	36	256	
Medium-Memory-1	35	36	768	
Medium-Memory-2	4	24	768	3 TB SSD local scratch
High-Memory-1	5	36	1,536	
High-Memory-2	6	40	1,536	3 TB SSD local scratch
GPGPU-1	36	36	256	NVIDIA P100 GPU has 3,584 CUDA cores and 16 GB CoWoS
MIC-1	36	64	384	16 GB of high bandwidth (400 GB/s) stacked memory
VDI	5	36	256	NVIDIA Quadro M5000 GPU
v100x8	3	36	768	8 NVIDIA V100 GPUs with 5,120 CUDA cores and 32 GB CoWoS
Faculty Partner Nodes	3			Various research specific NVIDIA GPU configurations
ManeFrame II	354	11,276	120 TB	2.8 PB storage and InfiniBand network



ManeFrame File Systems

- \$HOME
 - NFS file system, located in /users (e.g. /users/\$USER)
 - When you log in to ManeFrame your home directory is where you will land in by default
 - This space should be used to write, edit, compile, and browse your programs and job submission scripts
 - \$HOME space is restricted by quotas (200 GB)

ManeFrame II File Systems



\$WORK

- NFS file system, located at /work/users/\$USER
- \$WORK space is restricted by quotas (8 TB)
- Intended for long term storage



ManeFrame File Systems

- \$SCRATCH
 - 2.8 PB of scratch space is located in /scratch/users/ (e.g. /scratch/users/\$USER)
 - High performance parallel Lustre file system
 - As a high-performance parallel filesystem, Lustre will not perform well if misused
 - \$SCRATCH is a volatile file system
 - No guarantee that any of the files stored in \$SCRATCH can be retrieved or restored in the event of an accidental delete, loss or failure of the filesystem



University Data Center (UDC)



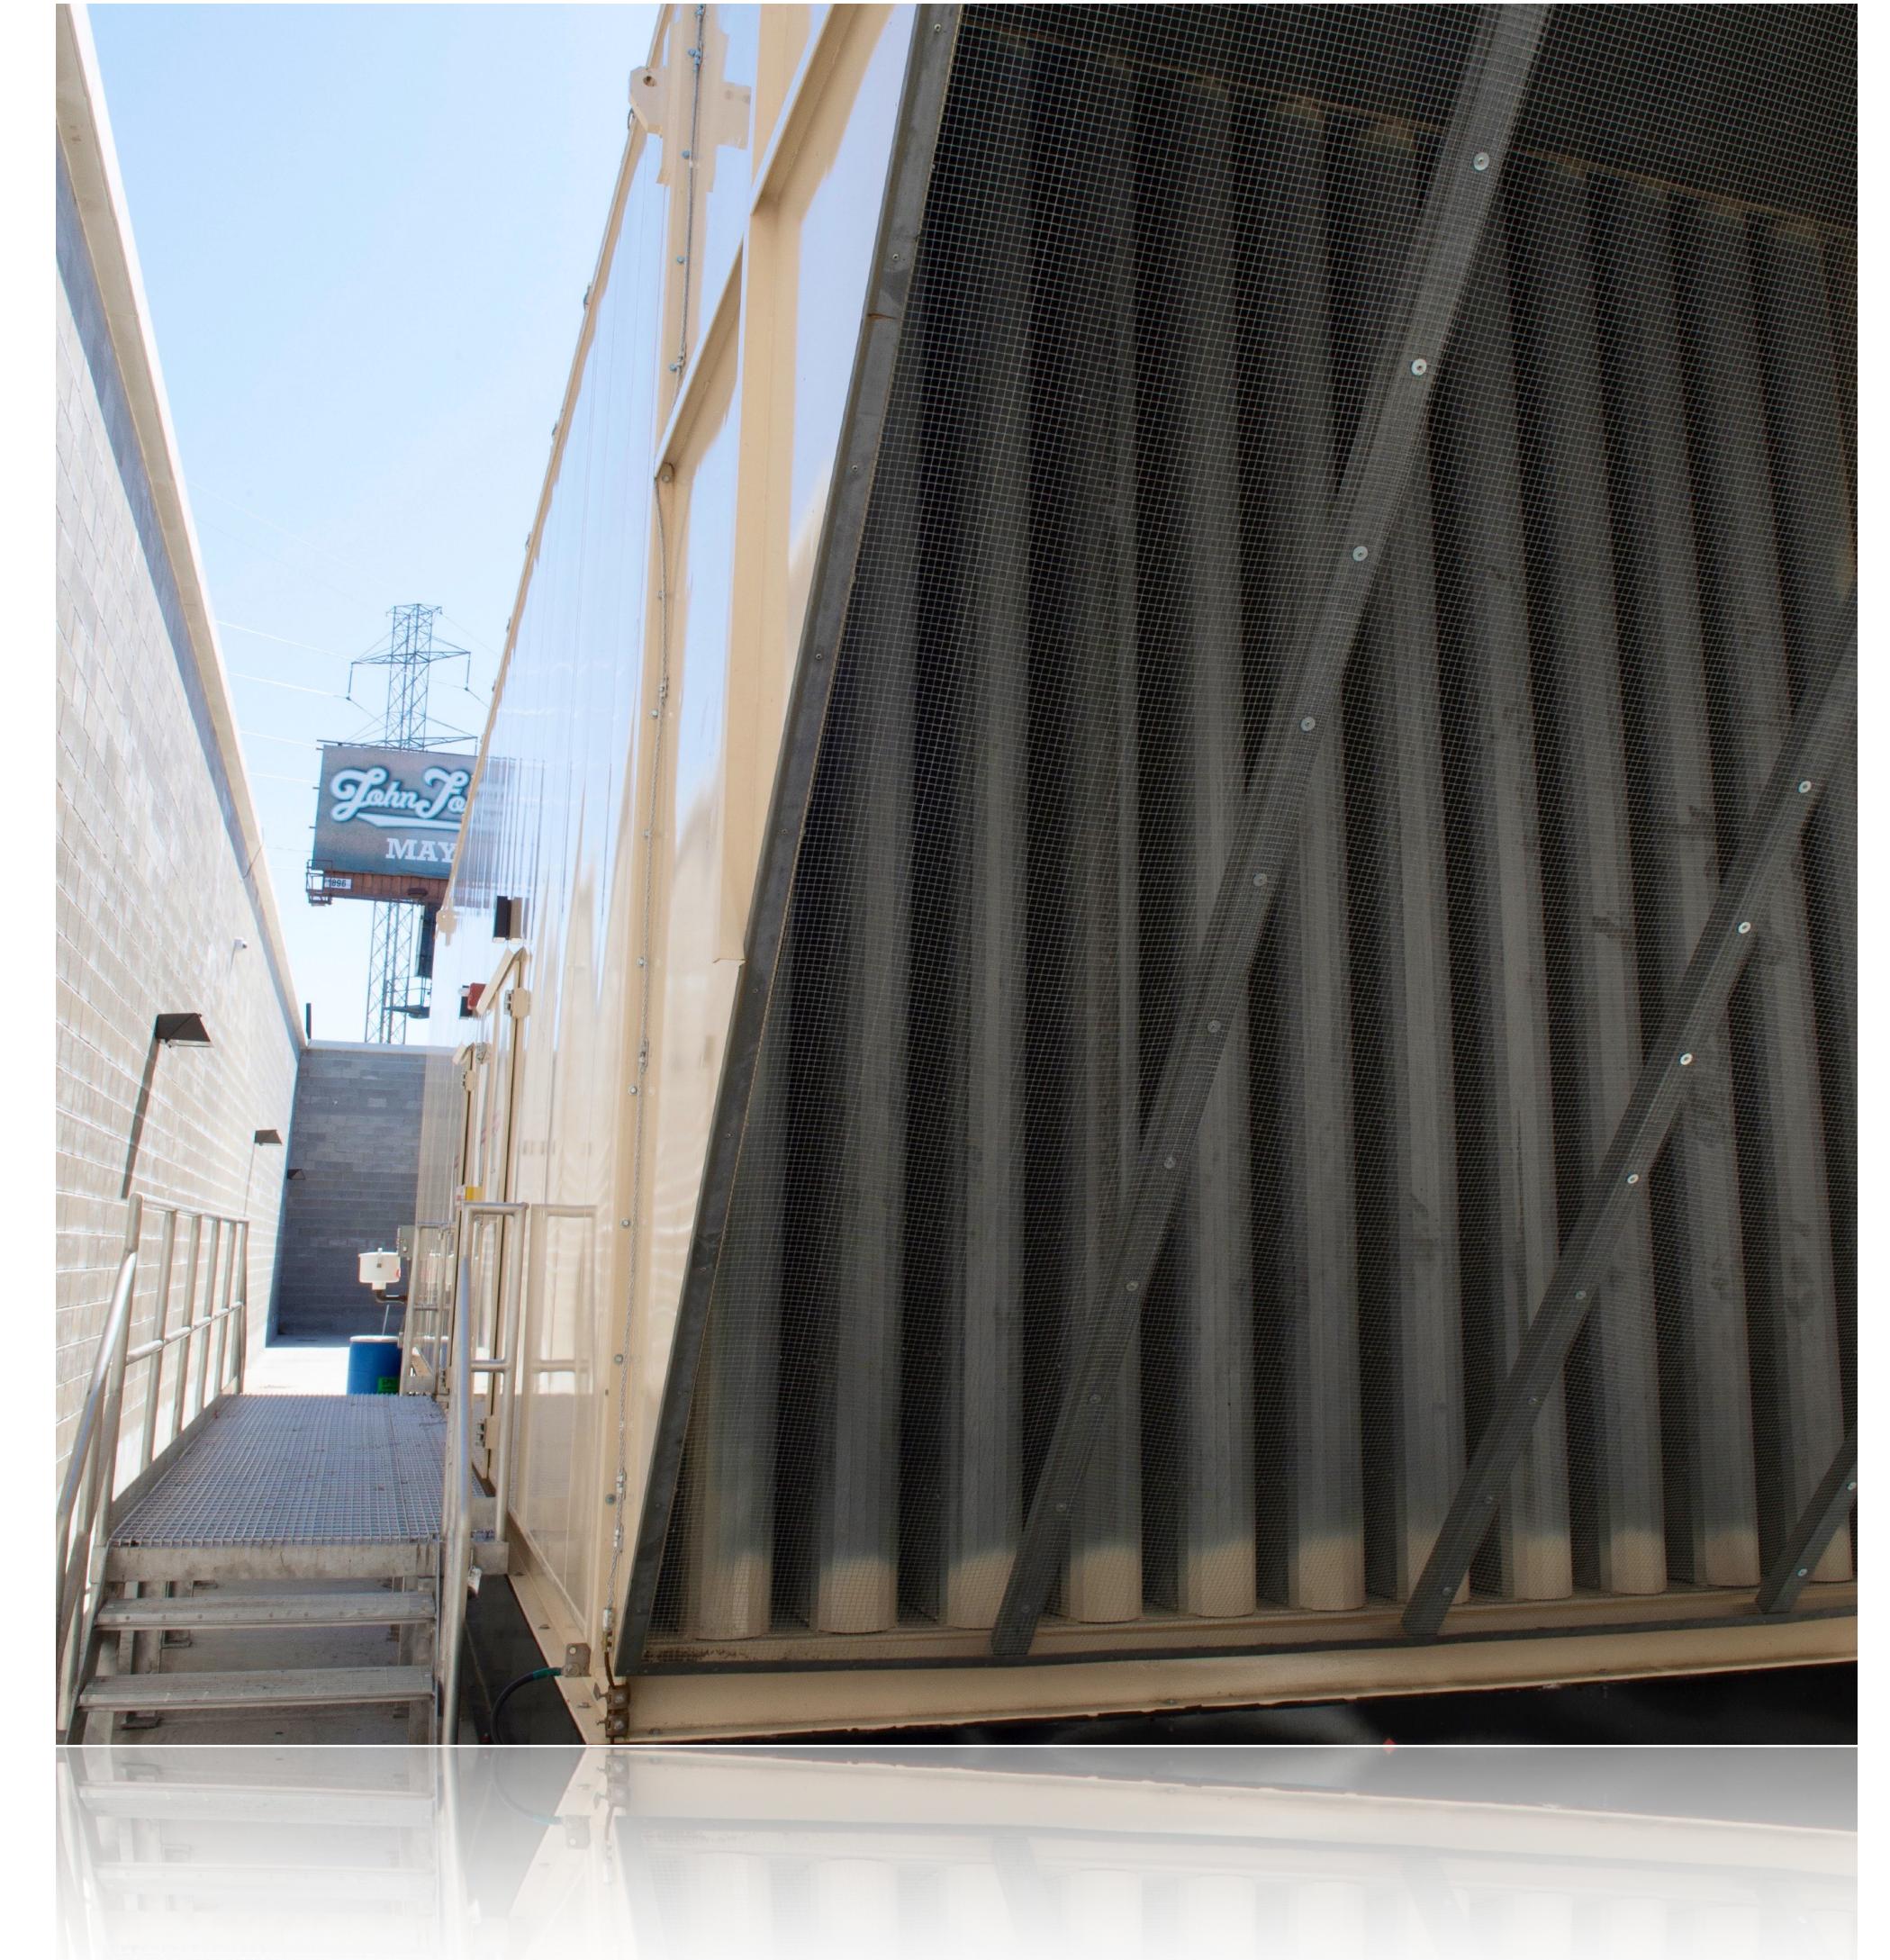
ManeFrame II



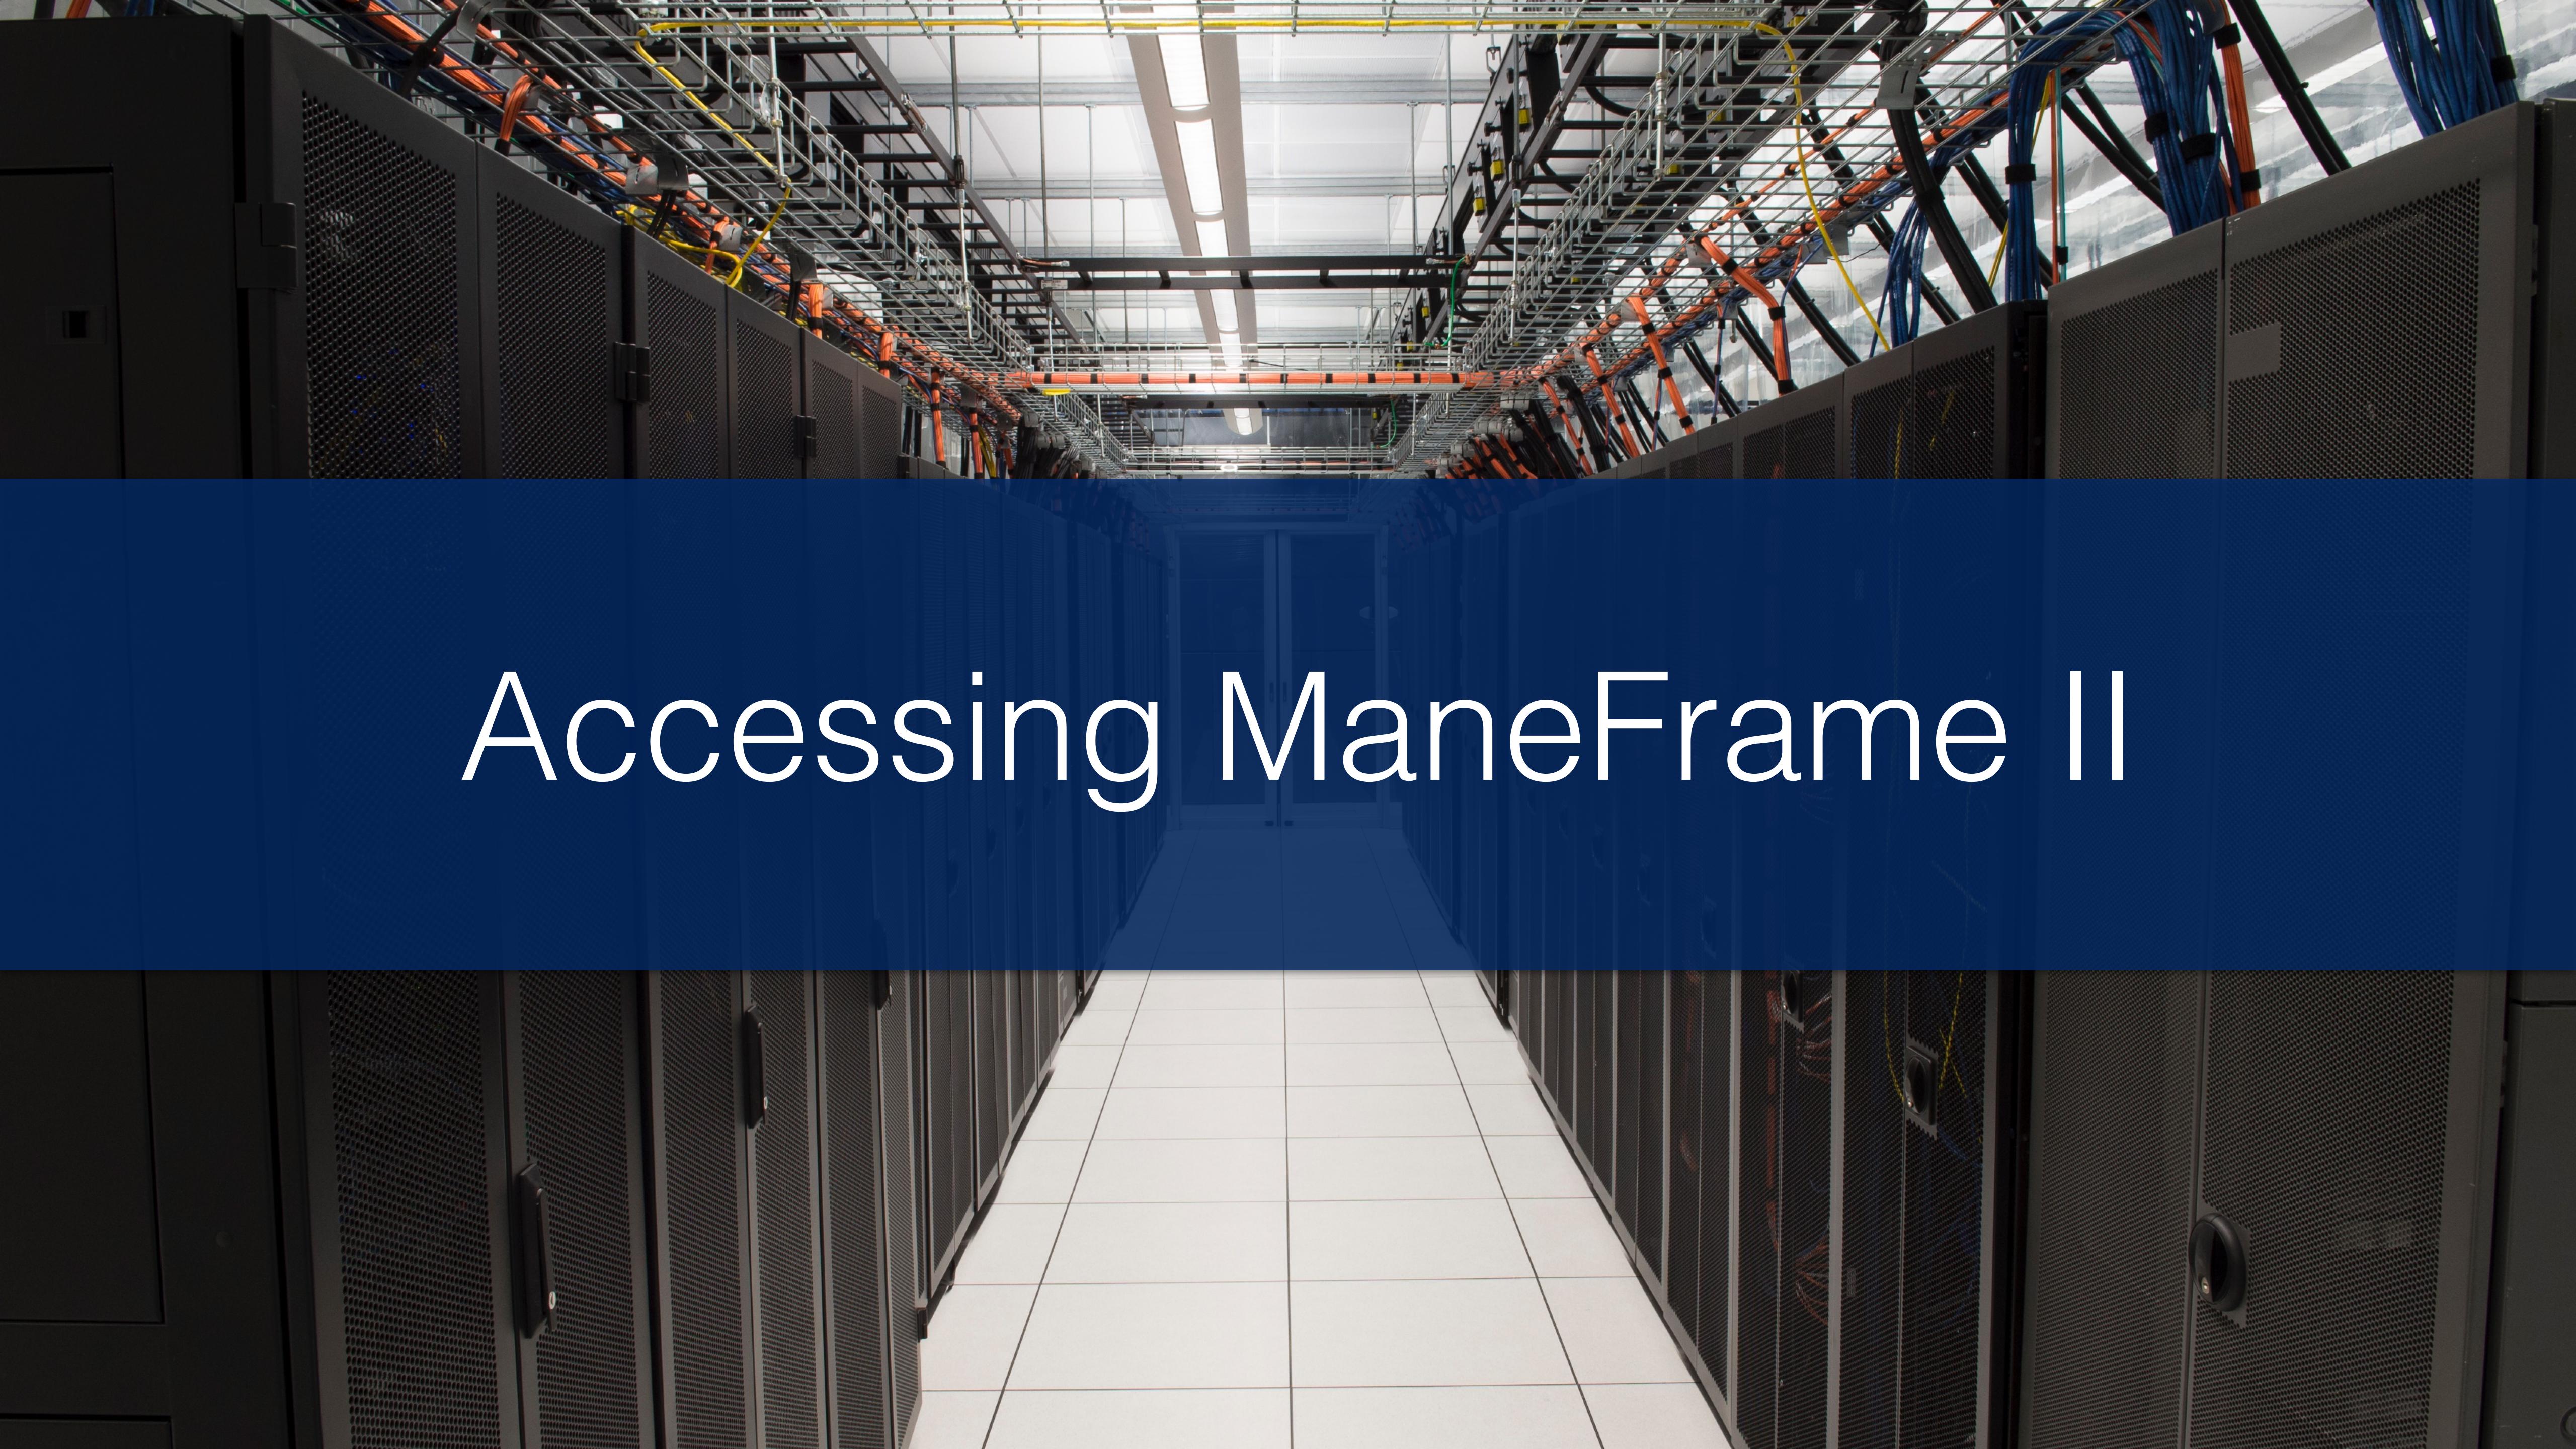
ManeFrame Hot Aisle



Cooling and Backup Generator



Accessing ManeFrame II





Accessing ManeFrame II

- SSH (secure shell)
 - Provides command line access to the login nodes
 - Graphical applications require X11-forwarding
 - M2 has five login nodes all accessible via m2.smu.edu
- HPC Portal
 - Provides web-based access to M2 resources
 - Web-based terminal access
 - JupyterLab (Notebooks)
 - RStudio
 - Remote Desktops



SSH & X11

- SSH
- Linux, macOS, Windows, WSL is installed by default
- Available via terminal applications
- X11
 - Linux
 - Installed by default
 - macOS
 - Apple provides X11 via the Xquartz open source project
 - <https://www.xquartz.org>
 - Windows
 - Xming
 - www.straightrunning.com/XmingNotes/



SSH and X11 for Windows

- As Windows is not a UNIX-like operating system standard utilities are not available by default
- Many SSH clients available with various features and styles
 - Putty
 - WinSCP
 - Tera Term
- X11 is available via Xming
- UNIX-like emulated environments
 - Ubuntu for Windows
 - Only on the most recent version of Windows 10
 - X11 not provided, but can use Xming
 - Cygwin
 - Can provide both SSH and X11



1. Open terminal application

2. At the command prompt, enter “ssh
username@m2.smu.edu -C -X”

- “username” is your ManeFrame username
 - “-C” requests compression, which can make the connection faster
 - “-X” requests X11 forwarding, which allows graphical programs to run from ManeFrame, but display on your computer



Xming and Putty Installation Notes

- Default installation options are fine for these installations

1. Install Xming

- Upon initial launch Windows firewall may ask for permission, select allow access

2. Install Xming-fonts

3. Install Putty

Releases

Website Releases	Version	State/Notes	Released	MD5 signatures	Size MB
Xming-Xming-x64	7.7.0.10	Website Release	14 Aug 2016	MD5 signatures	4.93 5.21
Xming-portablePuTTY-Xming-portablePuTTY-x64	7.7.0.10	Website Release	9 Aug 2016	MD5 signatures	1.82 1.90

See [Donations](#) for how to obtain a Donor Password.

Public Domain Releases	Version	State/Notes	Released	MD5 signature	Size MB
Xming-fonts	7.7.0.10	Public Domain	9 Aug 2016	ed1a0ab53688615bfec88ab399ae5470	31.1
Xming-Xming-mesa	6.9.0.31	Public Domain	4 May 2007	4cd12b9bec0ae19b95584650bbaf534a e580debbf6110cf4d8fcfd20beb541c1	2.10 2.50

See [Donations](#) for how to obtain a Donor Password.

Website Snapshots	Version	State/Notes	Snapshot	MD5 signature	Size MB
Snapshot Xming	7.7.0.11	Work in progress	7 Sep 2016 08:25	Not yet released	4.94 5.22

See [Donations](#) for how to obtain a Donor Password.

Binaries

The latest release version (beta 0.67)

This will generally be a version we think is reasonably likely to work well. If you have a problem, try an earlier release.

For Windows on Intel x86

PuTTY:	putty.exe	(or by FTP)	(signature)
PuTTYtel:	puttytel.exe	(or by FTP)	(signature)
PSCP:	pscp.exe	(or by FTP)	(signature)
PSFTP:	psftp.exe	(or by FTP)	(signature)
Plink:	plink.exe	(or by FTP)	(signature)
Pageant:	pageant.exe	(or by FTP)	(signature)
PuTTYgen:	puttygen.exe	(or by FTP)	(signature)

A ZIP file containing all the binaries (except PuTTYtel), and also the help files

Zip file: [putty.zip](#) (or by [FTP](#)) ([signature](#))

A Windows MSI installer package for everything except PuTTYtel

Installer: [putty-0.67-installer.msi](#) (or by [FTP](#)) ([signature](#))

Legacy Inno Setup installer. [Reportedly insecure!](#) Use with caution, if the MSI fails.

Legacy installer: [putty-0.67-installer.exe](#) (or by [FTP](#)) ([signature](#))

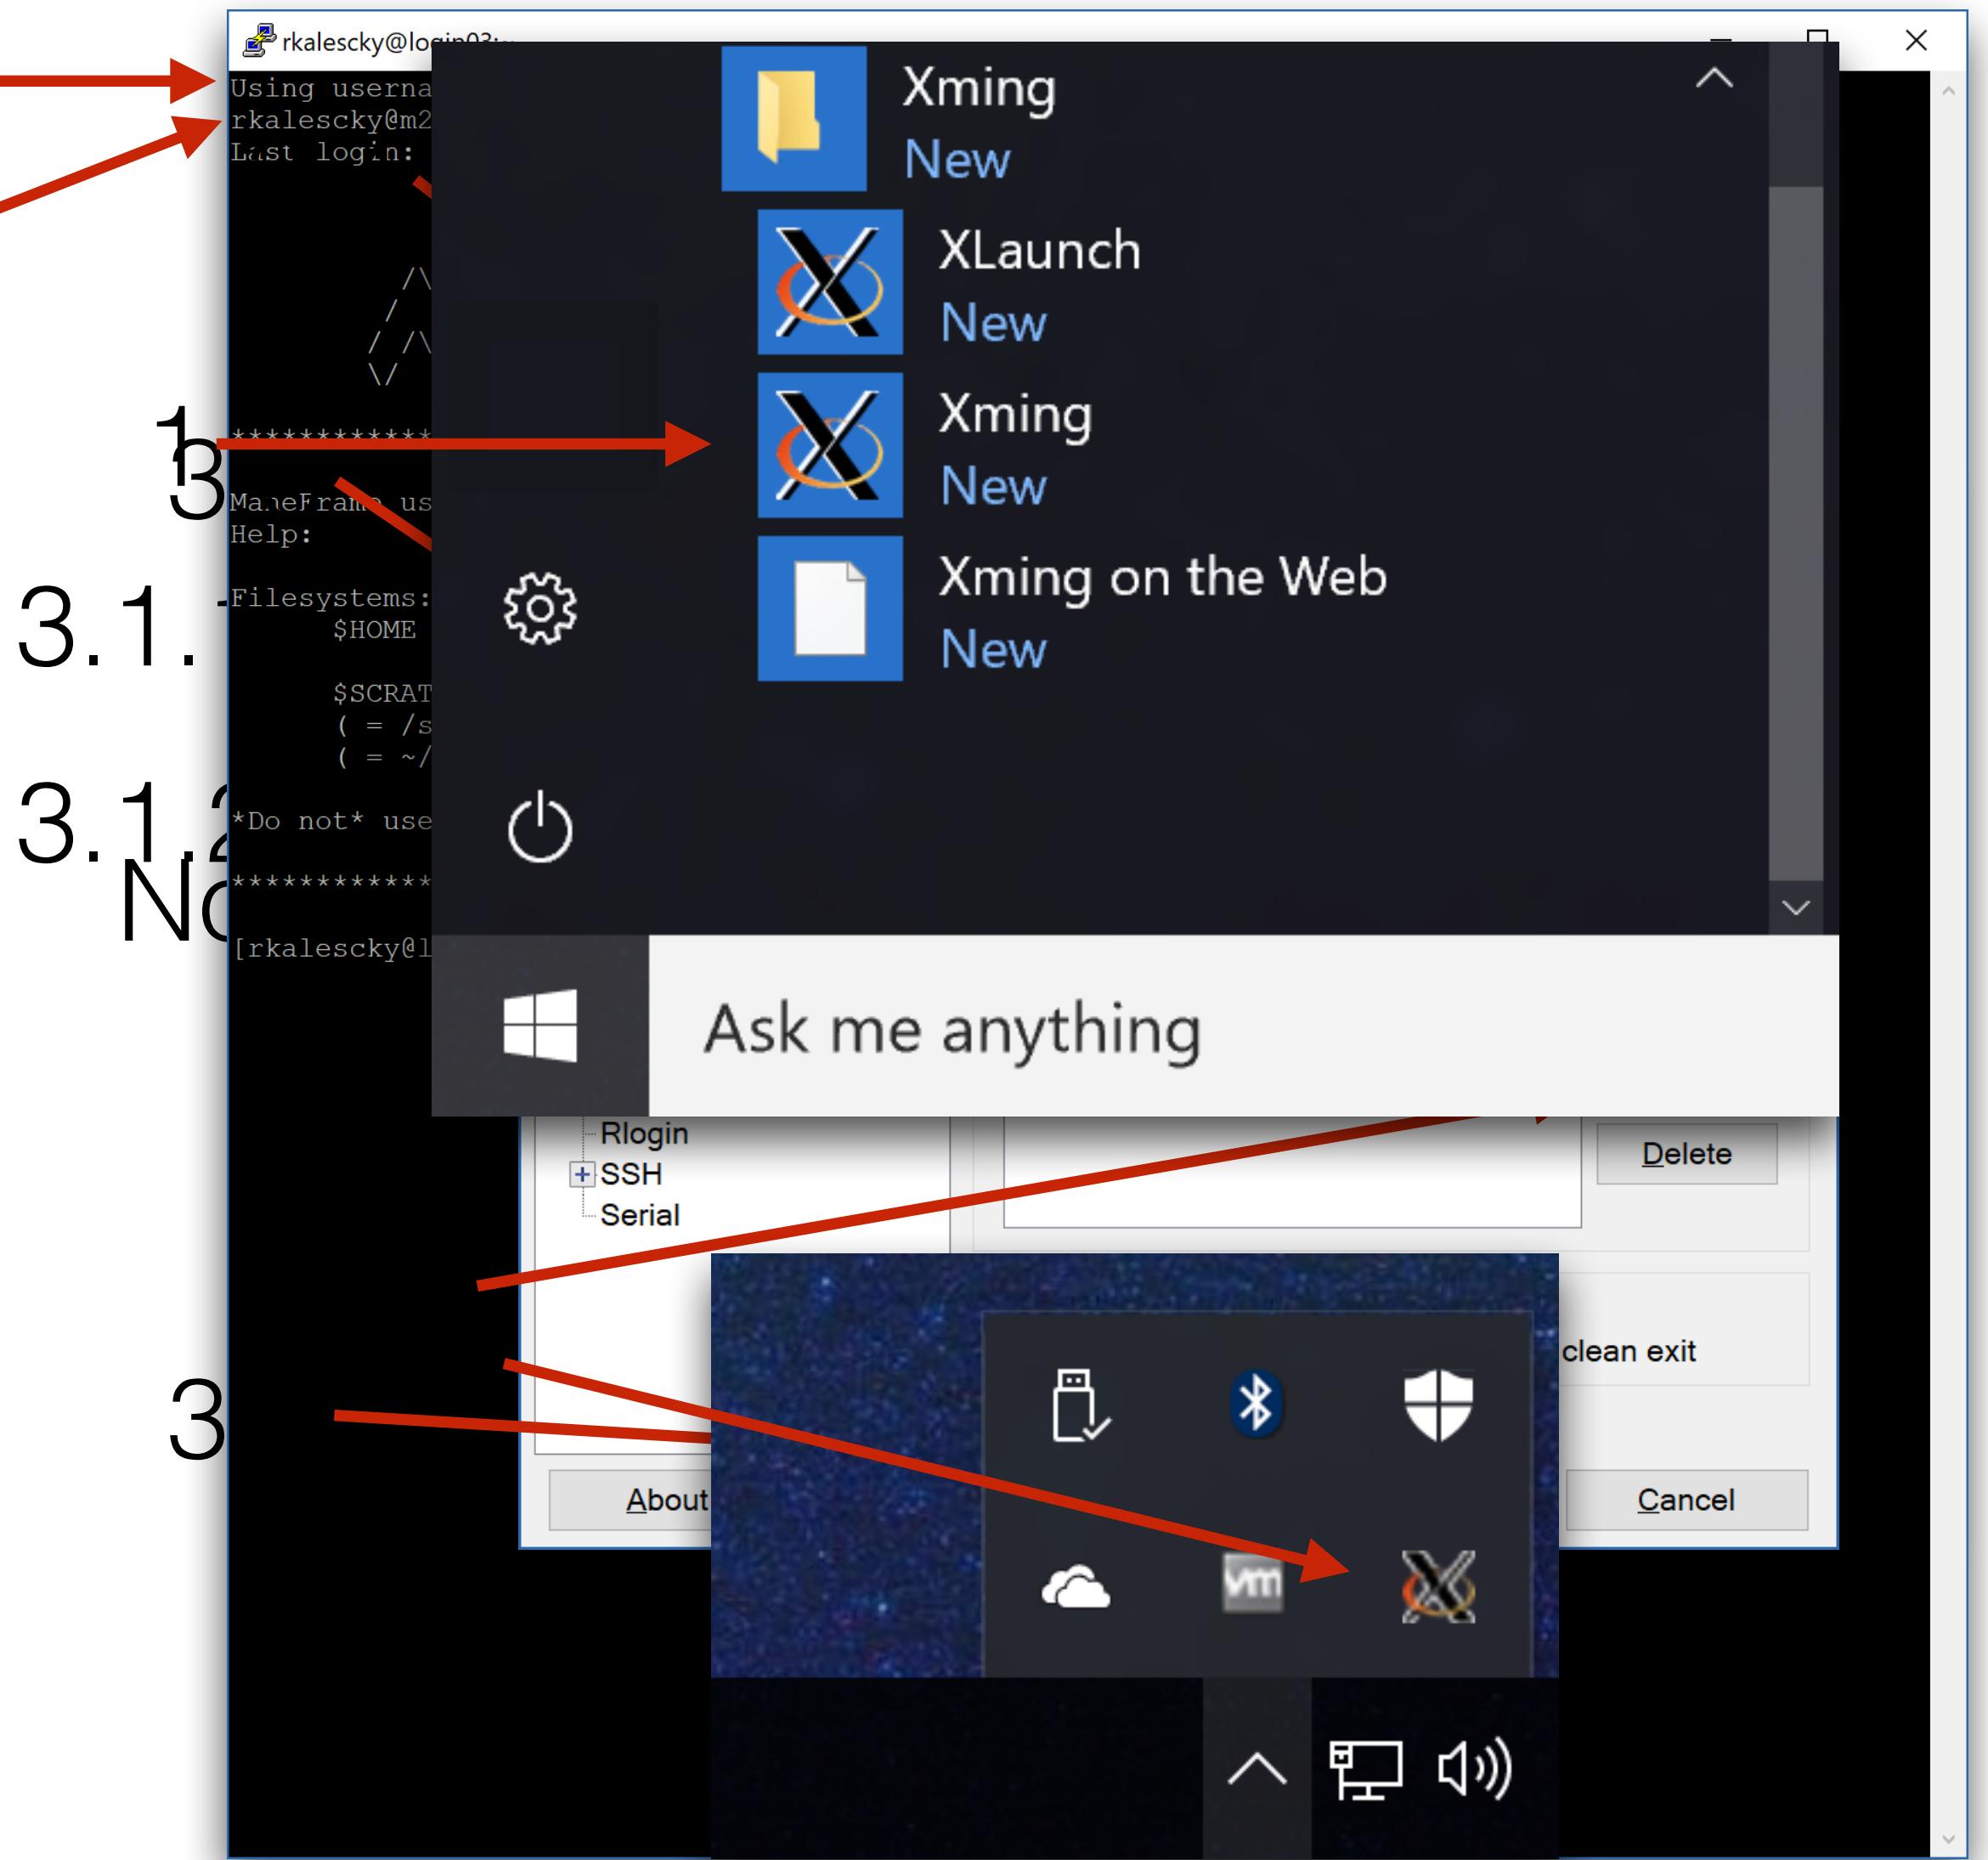
Checksums for all the above files

MD5:	md5sums	(or by FTP)	(signature)
SHA-1:	shalsums	(or by FTP)	(signature)
SHA-256:	sha256sums	(or by FTP)	(signature)
SHA-512:	sha512sums	(or by FTP)	(signature)

Accessing ManeFrame II: Putty



1. Start the Xming program in Start Menu
 - The Xming program runs from the System Tray
 2. Start the Putty program in the Start Menu
 3. In the Putty window:
 1. In the “Category:” box:
 1. Scroll to the bottom and select “+” next to “SSH”
 2. Select “X11”
 2. In the “X11 forwarding” section, select “Enable X11 forwarding”
 3. In the “Category:” box, scroll to the top and select “Session”
 4. In the “Host Name” field, type “m2.smu.edu”
 5. In the “Saved Sessions” field, type “ManeFrame II (M2)”
 6. Press “Save”
 7. Press “Open”
 - Select “Yes” if you presented with a “Putty Security Alert” window
 8. At the command prompt, type your ManeFrame username followed by enter
 9. At the command prompt, type your ManeFrame password followed by enter





Transferring Files

- It is easy to transfer files to and from ManeFrame II using an SFTP client
- Web
 - HPC Portal
- macOS
 - [Cyberduck](#)
 - [Transmit](#)
- Windows
 - [WinSCP](#)
 - [Cyberduck](#)
- UNIX-(like) systems also have command line options available

WinSCP Installation



Download WinSCP

WinSCP 5.9.1

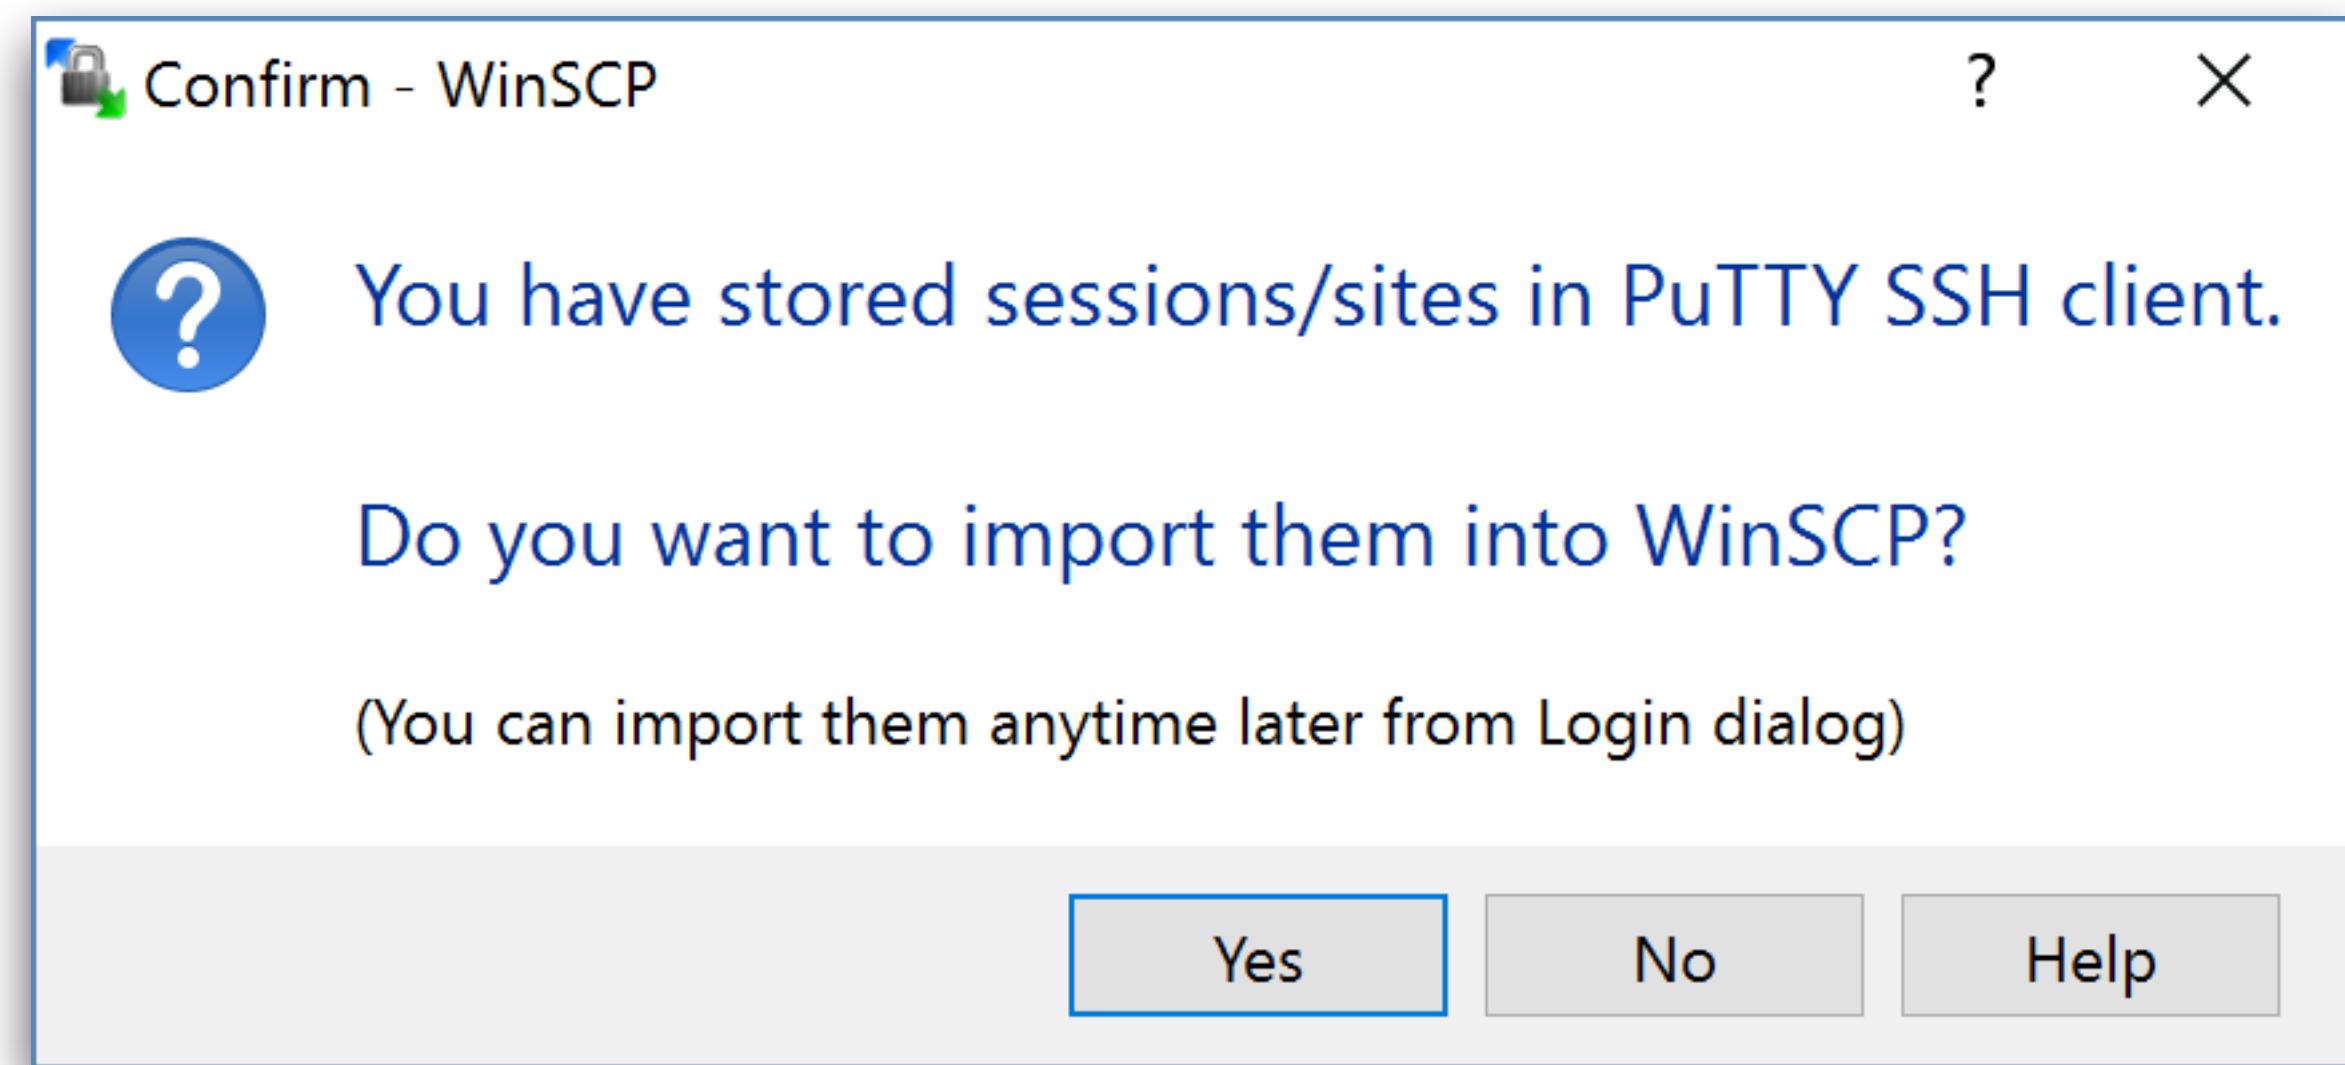
[Installation package](#) (8.6 MB; 533,570 downloads to date)

[Portable executables](#) (7.1 MB; 78,514 downloads to date)

[.NET assembly / COM library](#) (7.1 MB; 4,735 downloads to date)

[Source code](#) (11.9 MB; 1,734 downloads to date)

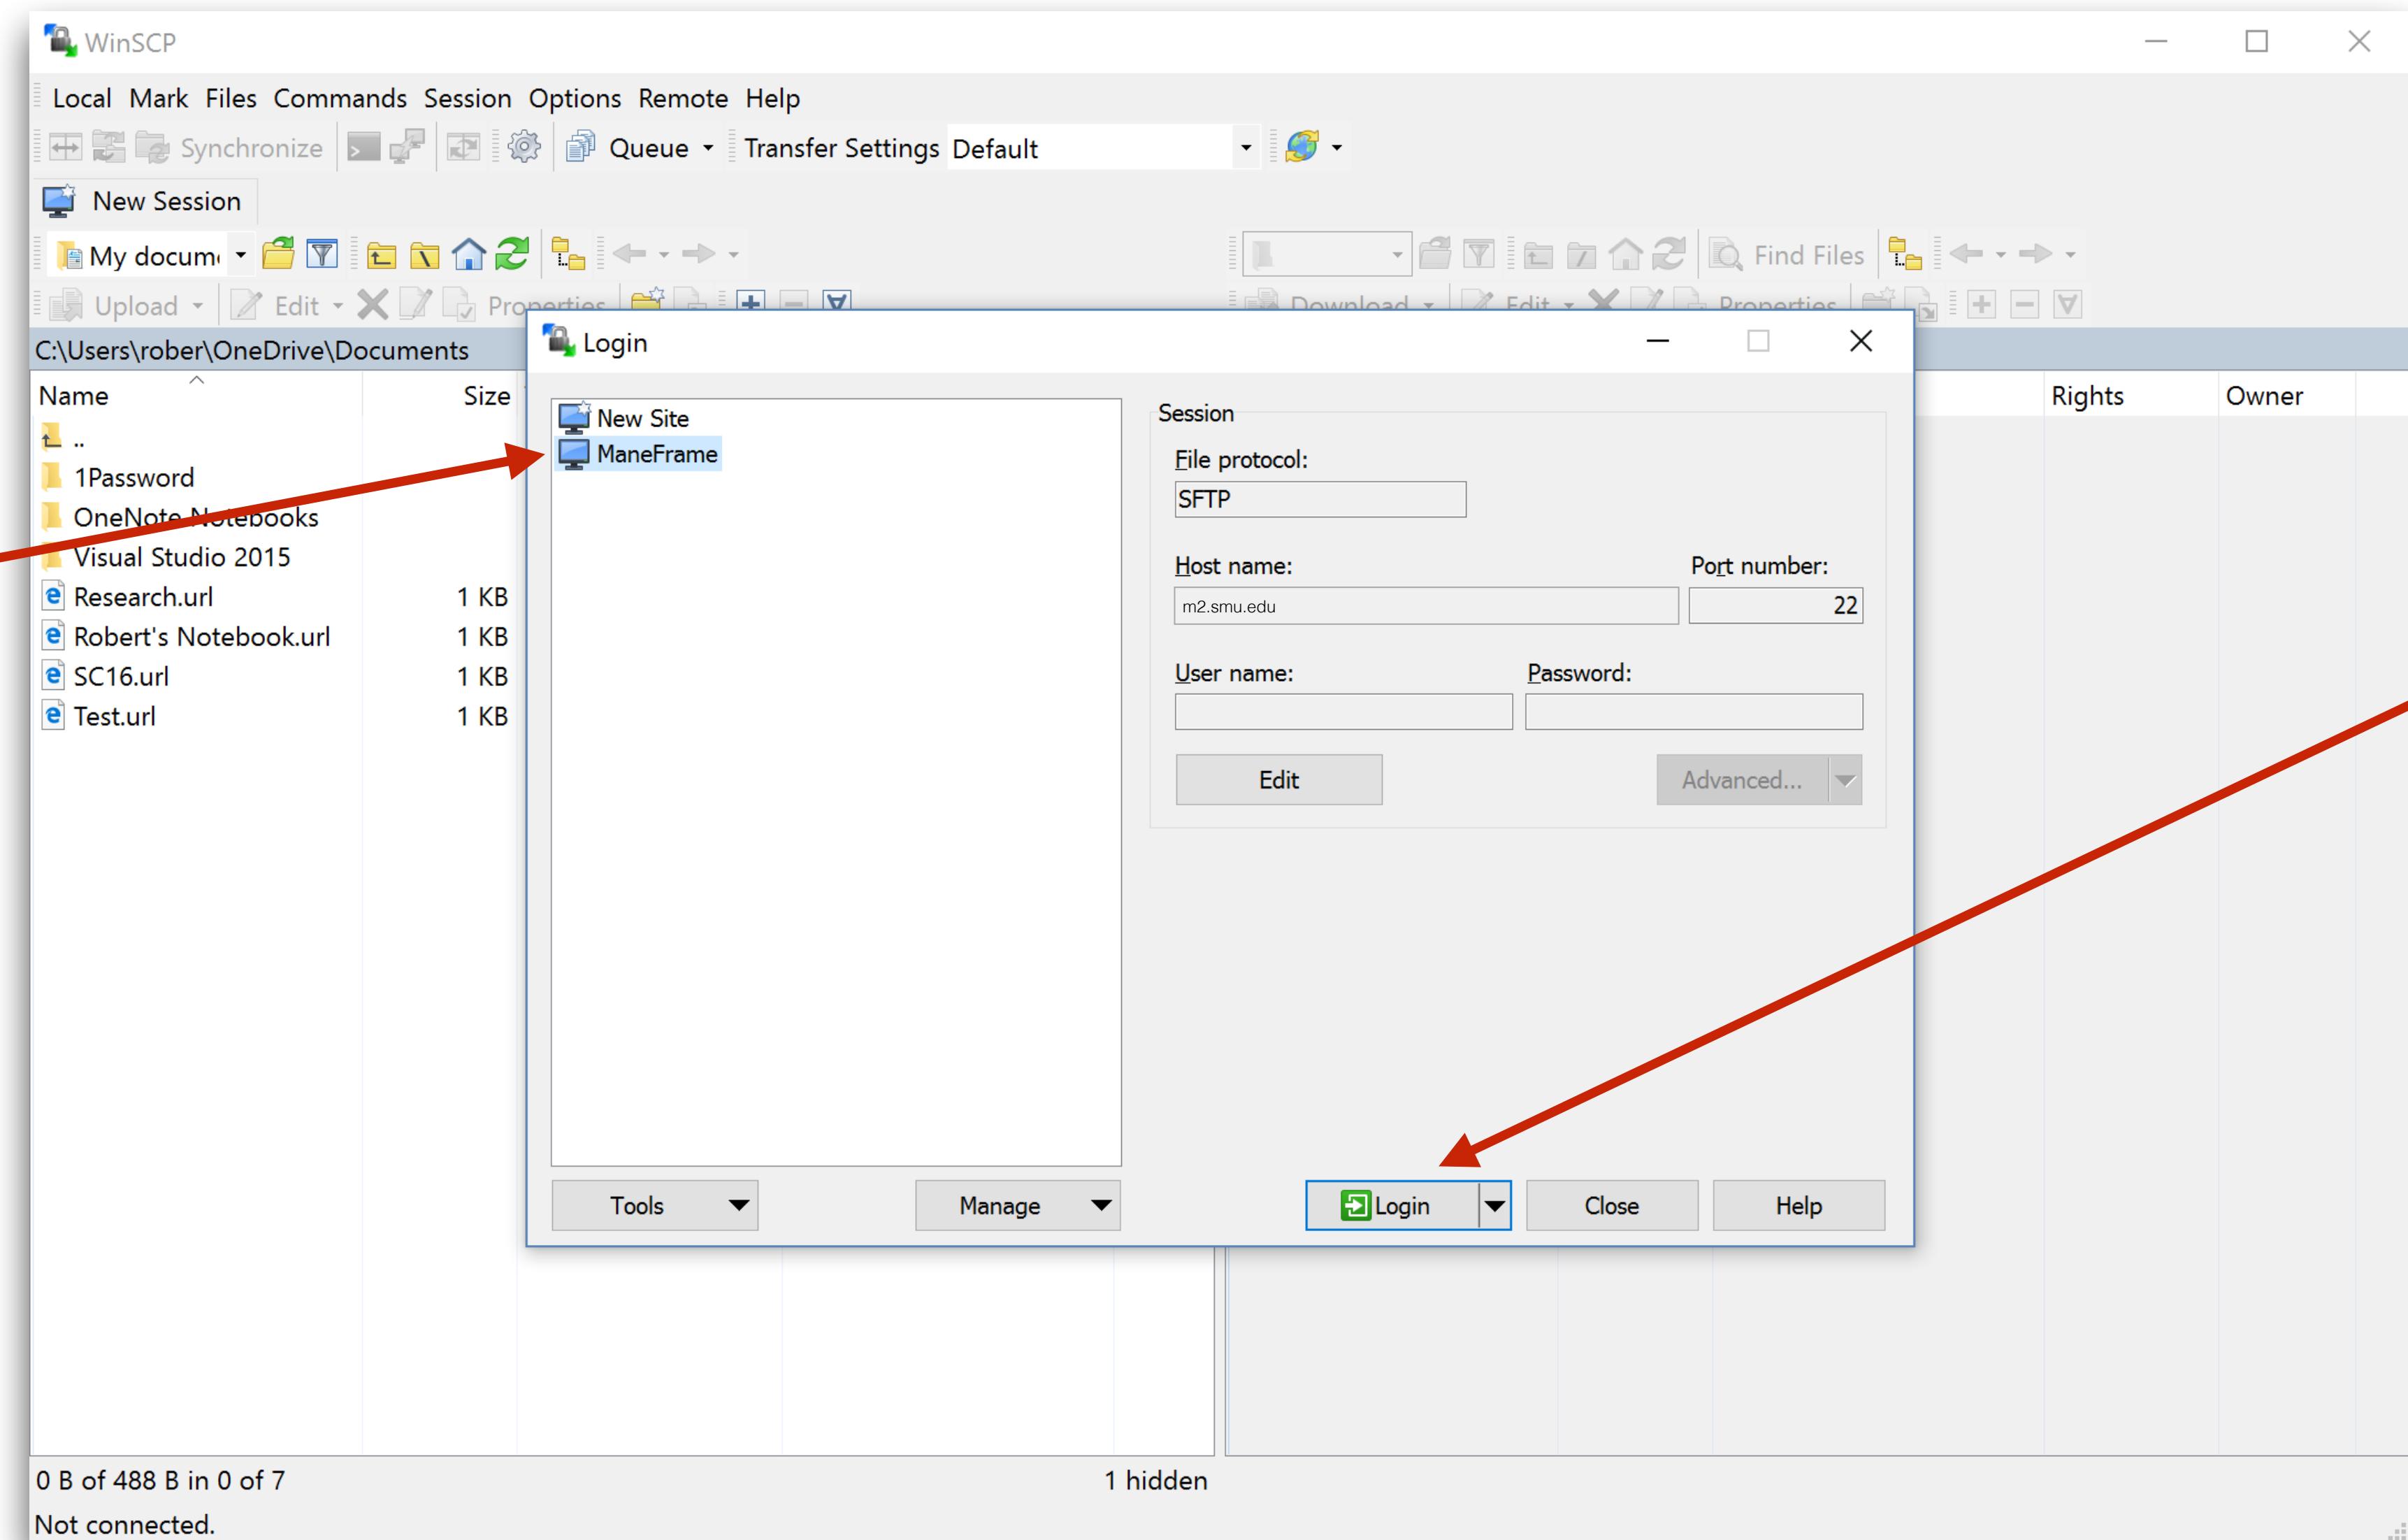
[\[Release Notes, Checksums\]](#) [\[What's New\]](#) [\[Release Notifications\]](#)





Using WinSCP

Can be imported
from Putty



You will be prompted
for your password



Using WinSCP

You can drag and drop across

Local Directories Remote Directories

Name	Size	Type	Changed	Rights	Owner
..			9/8/2016 11:01:58 AM	rwxr-xr-x	root
backups			9/8/2015 2:29:19 PM	rwx-----	rkalescky
bin			8/24/2016 2:46:22 PM	rwxr-x---	rkalescky
calculations			3/8/2016 1:43:04 PM	rwxrwxrwx	rkalescky
geant4_workdir			6/1/2016 12:26:44 PM	rwxr-xr-x	rkalescky
include			2/13/2015 11:26:40 AM	rwxr-x---	rkalescky
intel			2/18/2016 7:53:38 AM	rwxr-xr-x	rkalescky
lib			2/13/2015 11:24:57 AM	rwxr-x---	rkalescky
scripts			6/12/2015 7:55:44 AM	rwxr-x---	rkalescky
share			8/18/2015 7:12:26 AM	rwxr-x---	rkalescky
slurm			6/24/2015 1:00:54 PM	rwxr-x---	rkalescky
software			3/8/2016 1:43:17 PM	rwxrwxrwx	rkalescky
temp			3/8/2016 1:43:37 PM	rwxrwxrwx	rkalescky
testing			3/8/2016 1:40:04 PM	rwxrwxrwx	rkalescky
workshops			5/17/2016 7:02:39 AM	rwxrwxrwx	rkalescky
code_param	8 KB		5/19/2016 1:42:12 PM	rw-r--r--	rkalescky



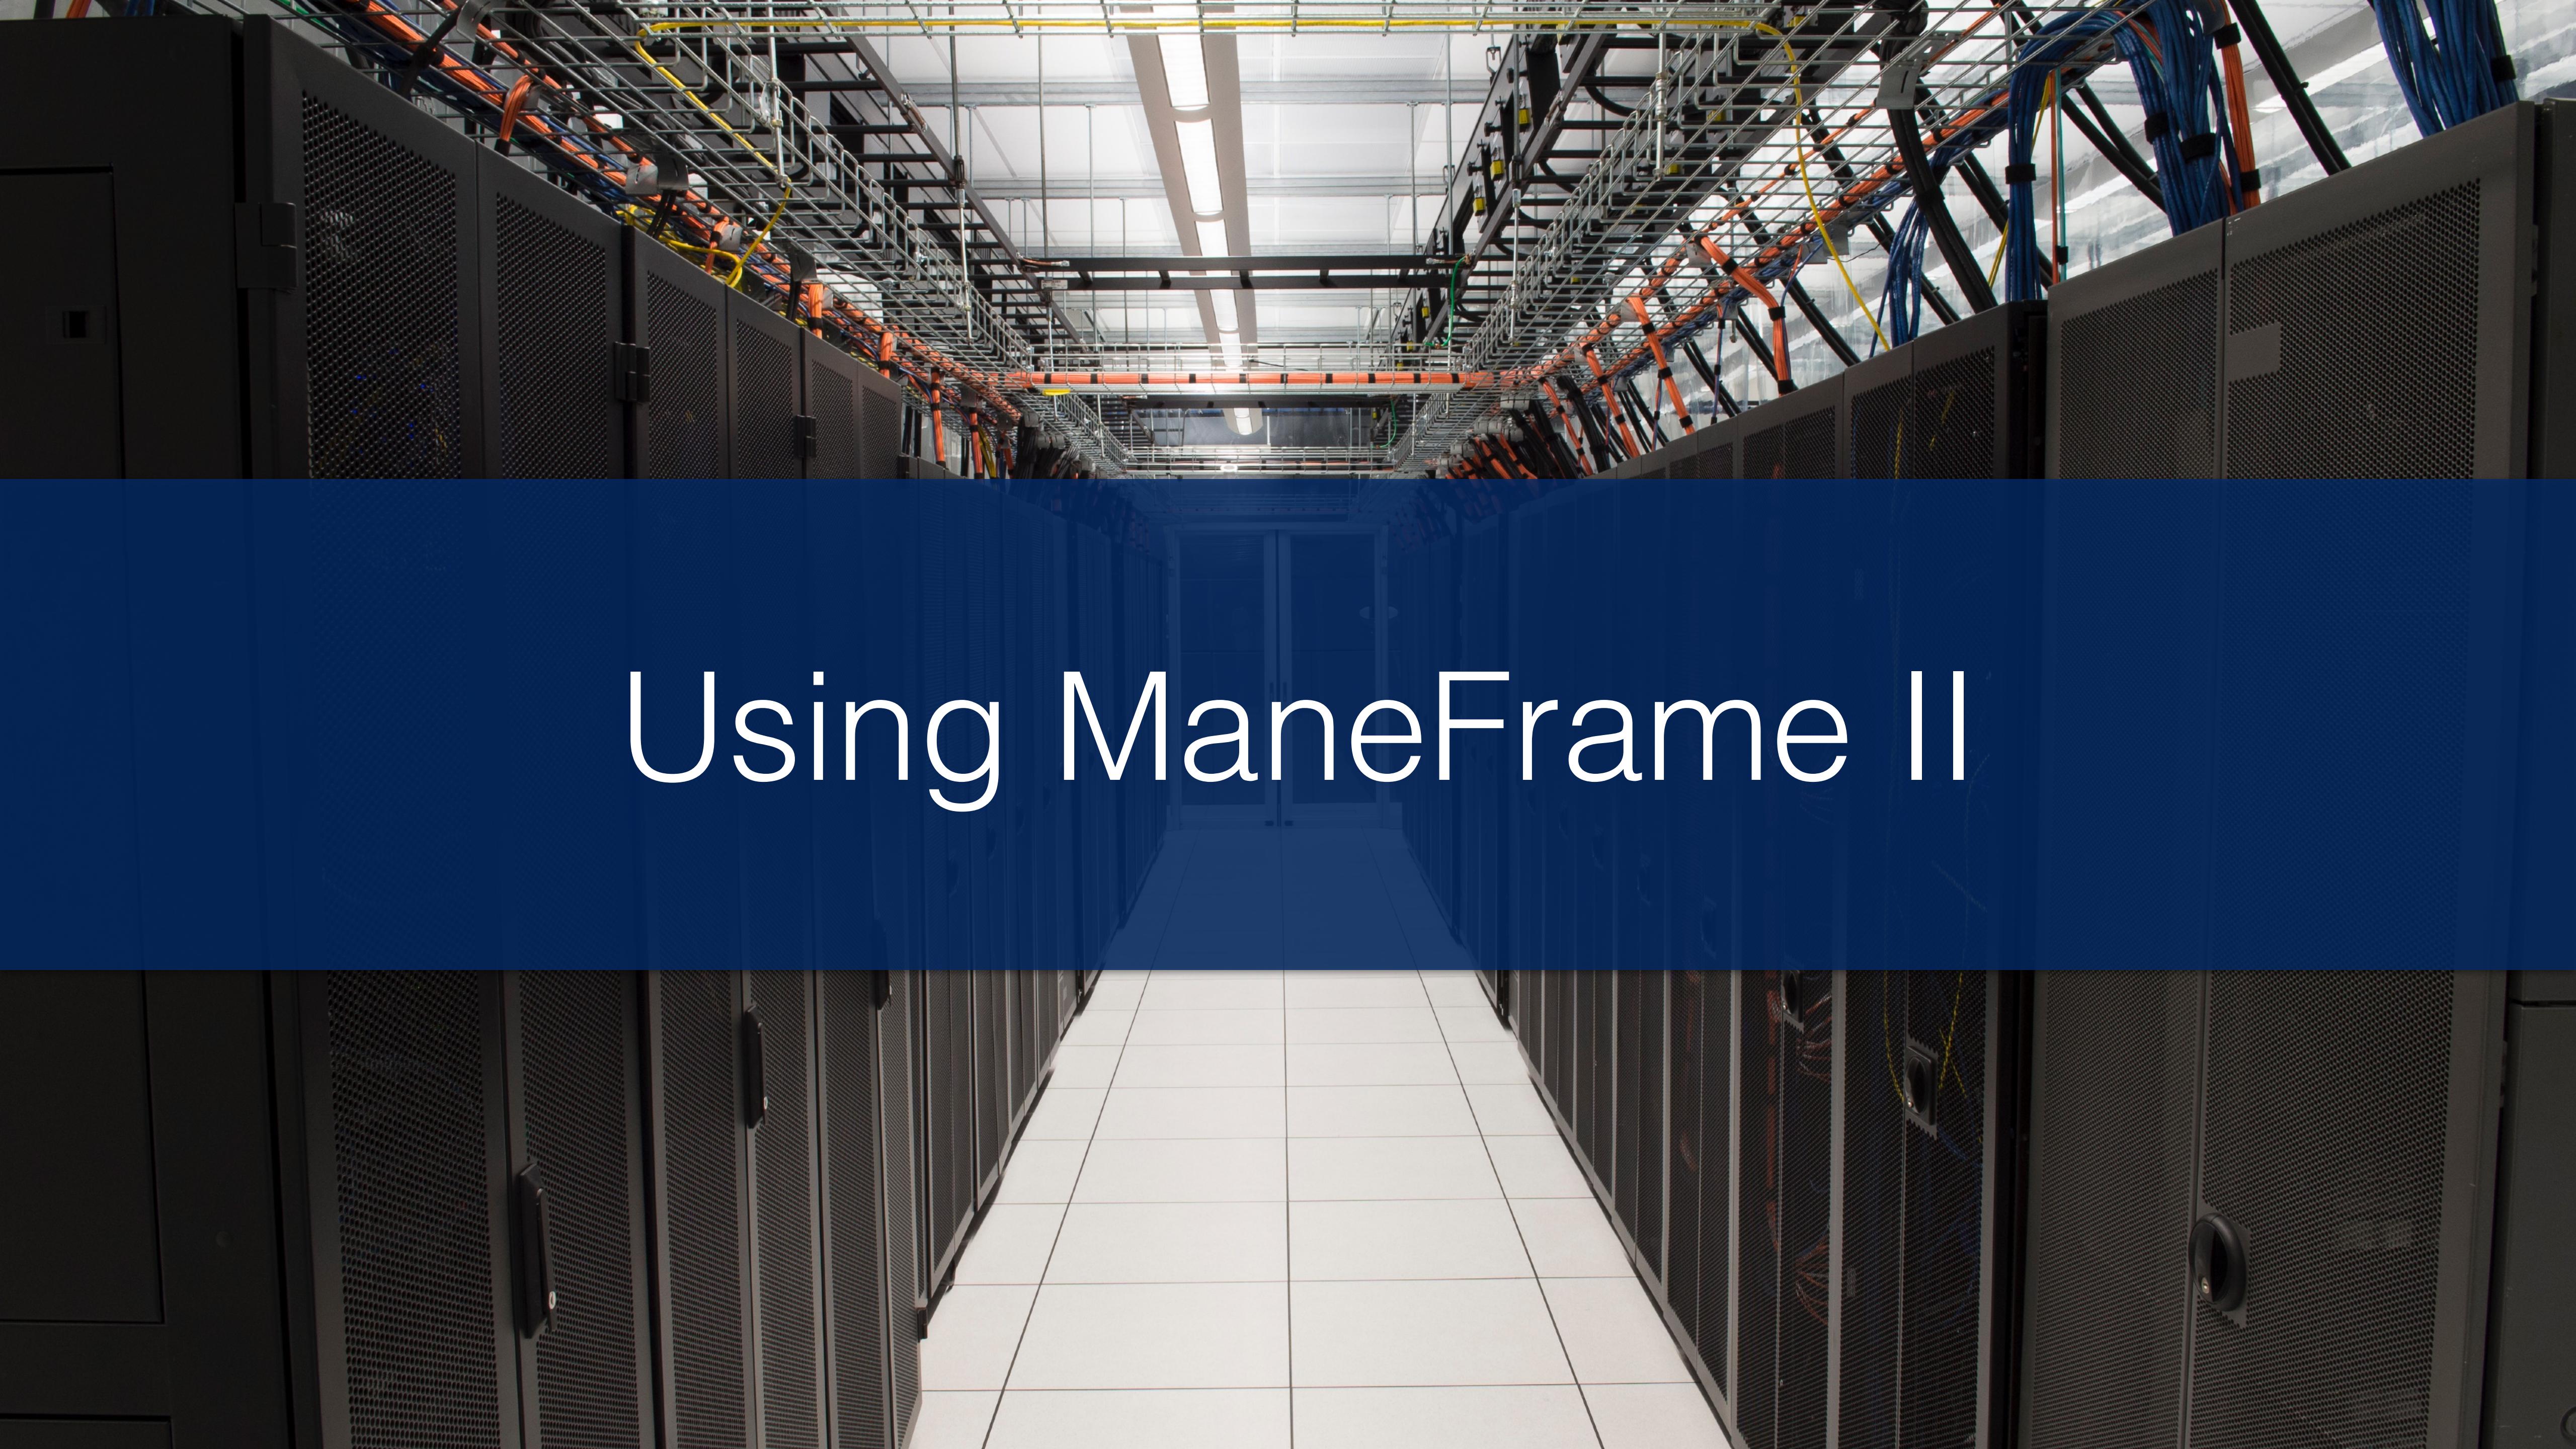
HPC Portal: hpc.smu.edu

- Provides an integrated, single access point for HPC resources on the ManeFrame II (M2) supercomputer.
- Accessing the Portal
 - Access to the HPC portal requires an existing M2 account.
 - Go to hpc.smu.edu.
 - Sign in using your SMU ID and SMU password
- Documentation
 - <http://s2.smu.edu/hpc/documentation/portal.html>

The screenshot shows the HPC OnDemand portal homepage. At the top, there's a navigation bar with links for 'HPC OnDemand', 'Files', 'Jobs', 'Clusters', 'Interactive Apps', and user information ('rkalesky' and 'Log Out'). Below the header, the SMU logo is displayed next to the text 'SMU'. To the right of the logo is a large image of a modern building with many windows, labeled 'ManeFrame II'. A sub-header below the logo reads: 'HPC OnDemand provides an integrated, single access point for HPC resources on the ManeFrame II supercomputer.' A section titled 'Message of the Day' features the heading 'Spring 2020 Center for Research Computing (CRC) Workshop Series'. It includes a paragraph about the workshop series, the date and time of the workshops (Tuesdays in Fondren Library East Room 110 from 2:00 to 4:00 PM), and a link to register. A table lists the workshop schedule:

Date	Workshop
January 21	M2 Introduction
January 28	Introduction to LAPACK and BLAS
February 4	Text Mining with Python on M2 (Lead by Dr. Eric Godat)
February 11	Using the New HPC Portal
February 18	Using GitHub

Using ManeFrame II





UNIX Tutorial

- UNIX tutorial
 - <http://s2.smu.edu/hpc/documentation/unix.html>

Accessing Applications and Tools



Software Available on M2



/hpc/modules/applications							
abcluster/2.0	compliance/3.0.2	geant4/10.06.p01	mesa/r12115	psi4/1.3.2-1	(D)	singularity/3.0.3	(D)
aimall/17.01.25	comsol/5.2a	gmt/5.4.3	molcas/8.2	python/2		singularity/3.5.3	
amber/16.15	(D) comsol/5.4	gnuplot/5.2.2	molpro/2015.1.24	python/3	(D)	snap/7.0.2	
amber/16/cpu	comsol/5.5	(D) gromacs/2019.4	molpro/2018.1.1	(D)	q-chem/5.0	snid/5.0	
amber/16/gpu	crystal/17/1.0.2	gurobi/8.0.1	mplus/8	q-chem/5.1	(D)	spack	
ansys/18.2/electronics	ctioga2/0.14.1	hdfview/3.0.0	mplus/8.2	(D)	q-chem/5.2.1	star/2.5.3a	
ansys/19.1/electronics	cuda/10.0	idl/8.0	namd/2.12/cpu	(D)	quantum-espresso/6.4.1	stata/mp-14	
ansys/19.1/workbench	(D) cudnn/7.6.4	jags/4.3.0	namd/2.12/gpu		r/3.4.0	stata/mp-15	(D)
arrayfire/3.5.0	demon/6.0.2	julia/0.5.2	namd/2.13_cpu		r/3.4.1	tcad/2019	
autodock_vina/1.1.2	demon_nano/2019-06-24	julia/0.6.0	namd/2.13_gpu		r/3.4.1-1	tensorflow/1.2_gpu	
awscli/2	ds9/7.5	julia/1.0.1	(D) nbo/6.0		r/3.4.4	tensorflow/1.4_gpu	
babel/2.4.1	firefox/76	knitro/11.1.0	nccl/2.4.8		r/3.5.2	tensorflow/1.10_gpu	
bazel/0.26.1	gamess/2017r1	lammps/aug17	omnia/1		r/3.6.2	tensorflow/2.0_gpu	(D)
bazel/0.29.1	(D) gams/25.1.2	lammps/2016.11.17	(D) openmm/7.2		rstudio/1.1.383	tensorrt/5.1.5	
bedtools2/2.26.0	gauss/17.0.3	macaulay2/1.11	orca/4.0.0		rstudio/1.1.423	texlive/2017	
borg/1.8	gaussian/g03e	madgraph/2.6.3.2	orca/4.0.1	(D)	rstudio/1.1.442	(D)	valgrind/3.12.0
cfour/2.00beta	gaussian/g09d	mallet/2.0.8	orca/4.1.1		samtools/1.6	vasp/5.4.4/cpu	
cfour/2.1	(D) gaussian/g16a	mathematica/11.1.1	paraview/5.4.1		sapt/2016.1	vasp/5.4.4/gpu	(D)
charmm/c41b1	gaussian/g16b	matlab/r2016b	period04/1.2.0		sas/9.4.update6	vmd/1.9.3/egl	
clocksim/2014	gaussian/g16c	(D) matlab/r2017a	periodic_nbo/1.1		sas/9.4	vmd/1.9.3/ogl	(D)
cmake/3.11.1	gaussview/6.0.16	matlab/r2018b	(D) psi4/1.3.2		sextractor/2.19.5		
/hpc/modules/compilers							
gcc-4.8.5	gcc-7.3	intel-17.0.4	intel-2020.0	pgi-12.10	pgi-16.5	pgi-17.7	pgi-18.4
gcc-6.3	gcc-9.2	intel-18.0.1	llvm-4.0.0	pgi-13.2	pgi-17.4	pgi-18.10	pgi-19.5



Basic Module Commands

Command	Description	Example Usage
module avail	List available modules	module avail
module load	Load modules	module load python
module list	List loaded modules	module list
module unload	Unload modules	module unload gcc
module purge	Unload all modules	module purge
module spider	Search for modules matching text	module spider python
module save	Save loaded modules	module save



Module Tutorial

[http://s2.smu.edu/hpc/documentation/
modules.html](http://s2.smu.edu/hpc/documentation/modules.html)

Submitting Jobs Using Slurm





Common Classes of Applications

- Serial
 - Users run short-lived single-threaded calculations in batches
 - Not uncommon to see batches of greater than 10K
 - Batches can be submitted individually, via SLURM's array feature, or bundled to use a whole node
- Shared-Memory
 - Users run long-lived computationally intensive jobs with various memory requirements
 - Not uncommon to see run times of several months
- Distributed-Memory
 - Users run computationally intensive jobs that span nodes
 - Not uncommon to see hundreds of nodes used for single job



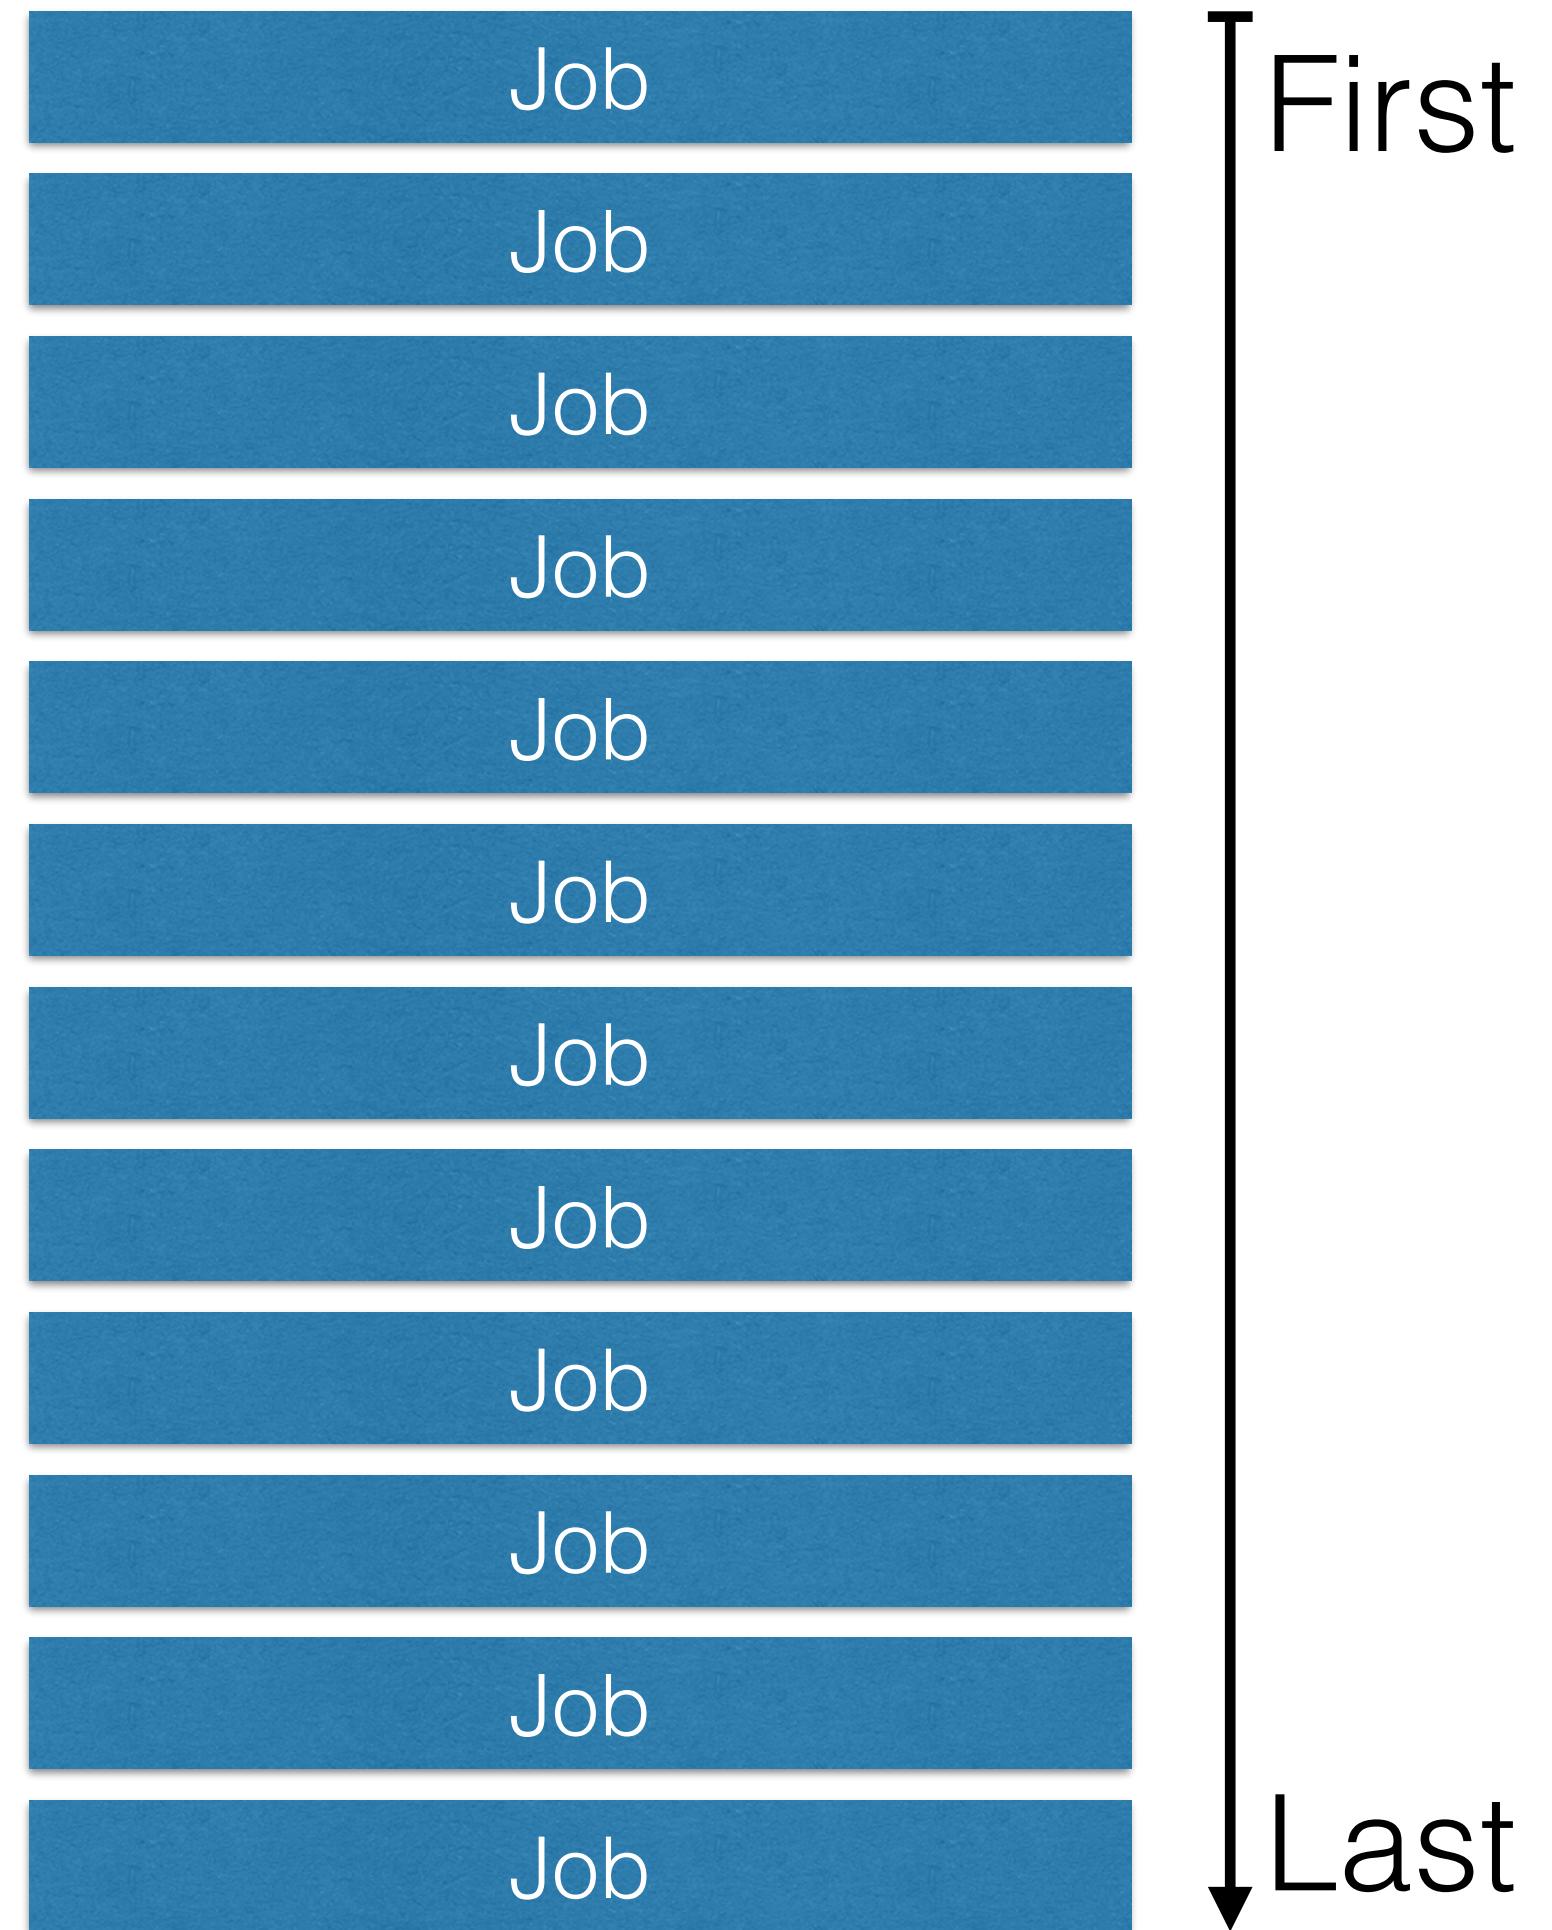
Common Classes of Applications

- Serial
 - Simple calculations frequently done *en masse*
 - Usually done to get statistics over many trials
 - R, MATLAB, Python, ROOT, Autodock
- Shared-Memory
 - Complex calculations with demanding compute requirements
 - Quantum chemistry and molecular dynamics packages (e.g. CFOUR, Gaussian, NAMD, CHARMM)
- Distributed-Memory
 - Distributed programs usually via MPI
 - Many parallel computing research codes
 - Some packages such as CFOUR, CHARMM, LAMMPS



Resource Management

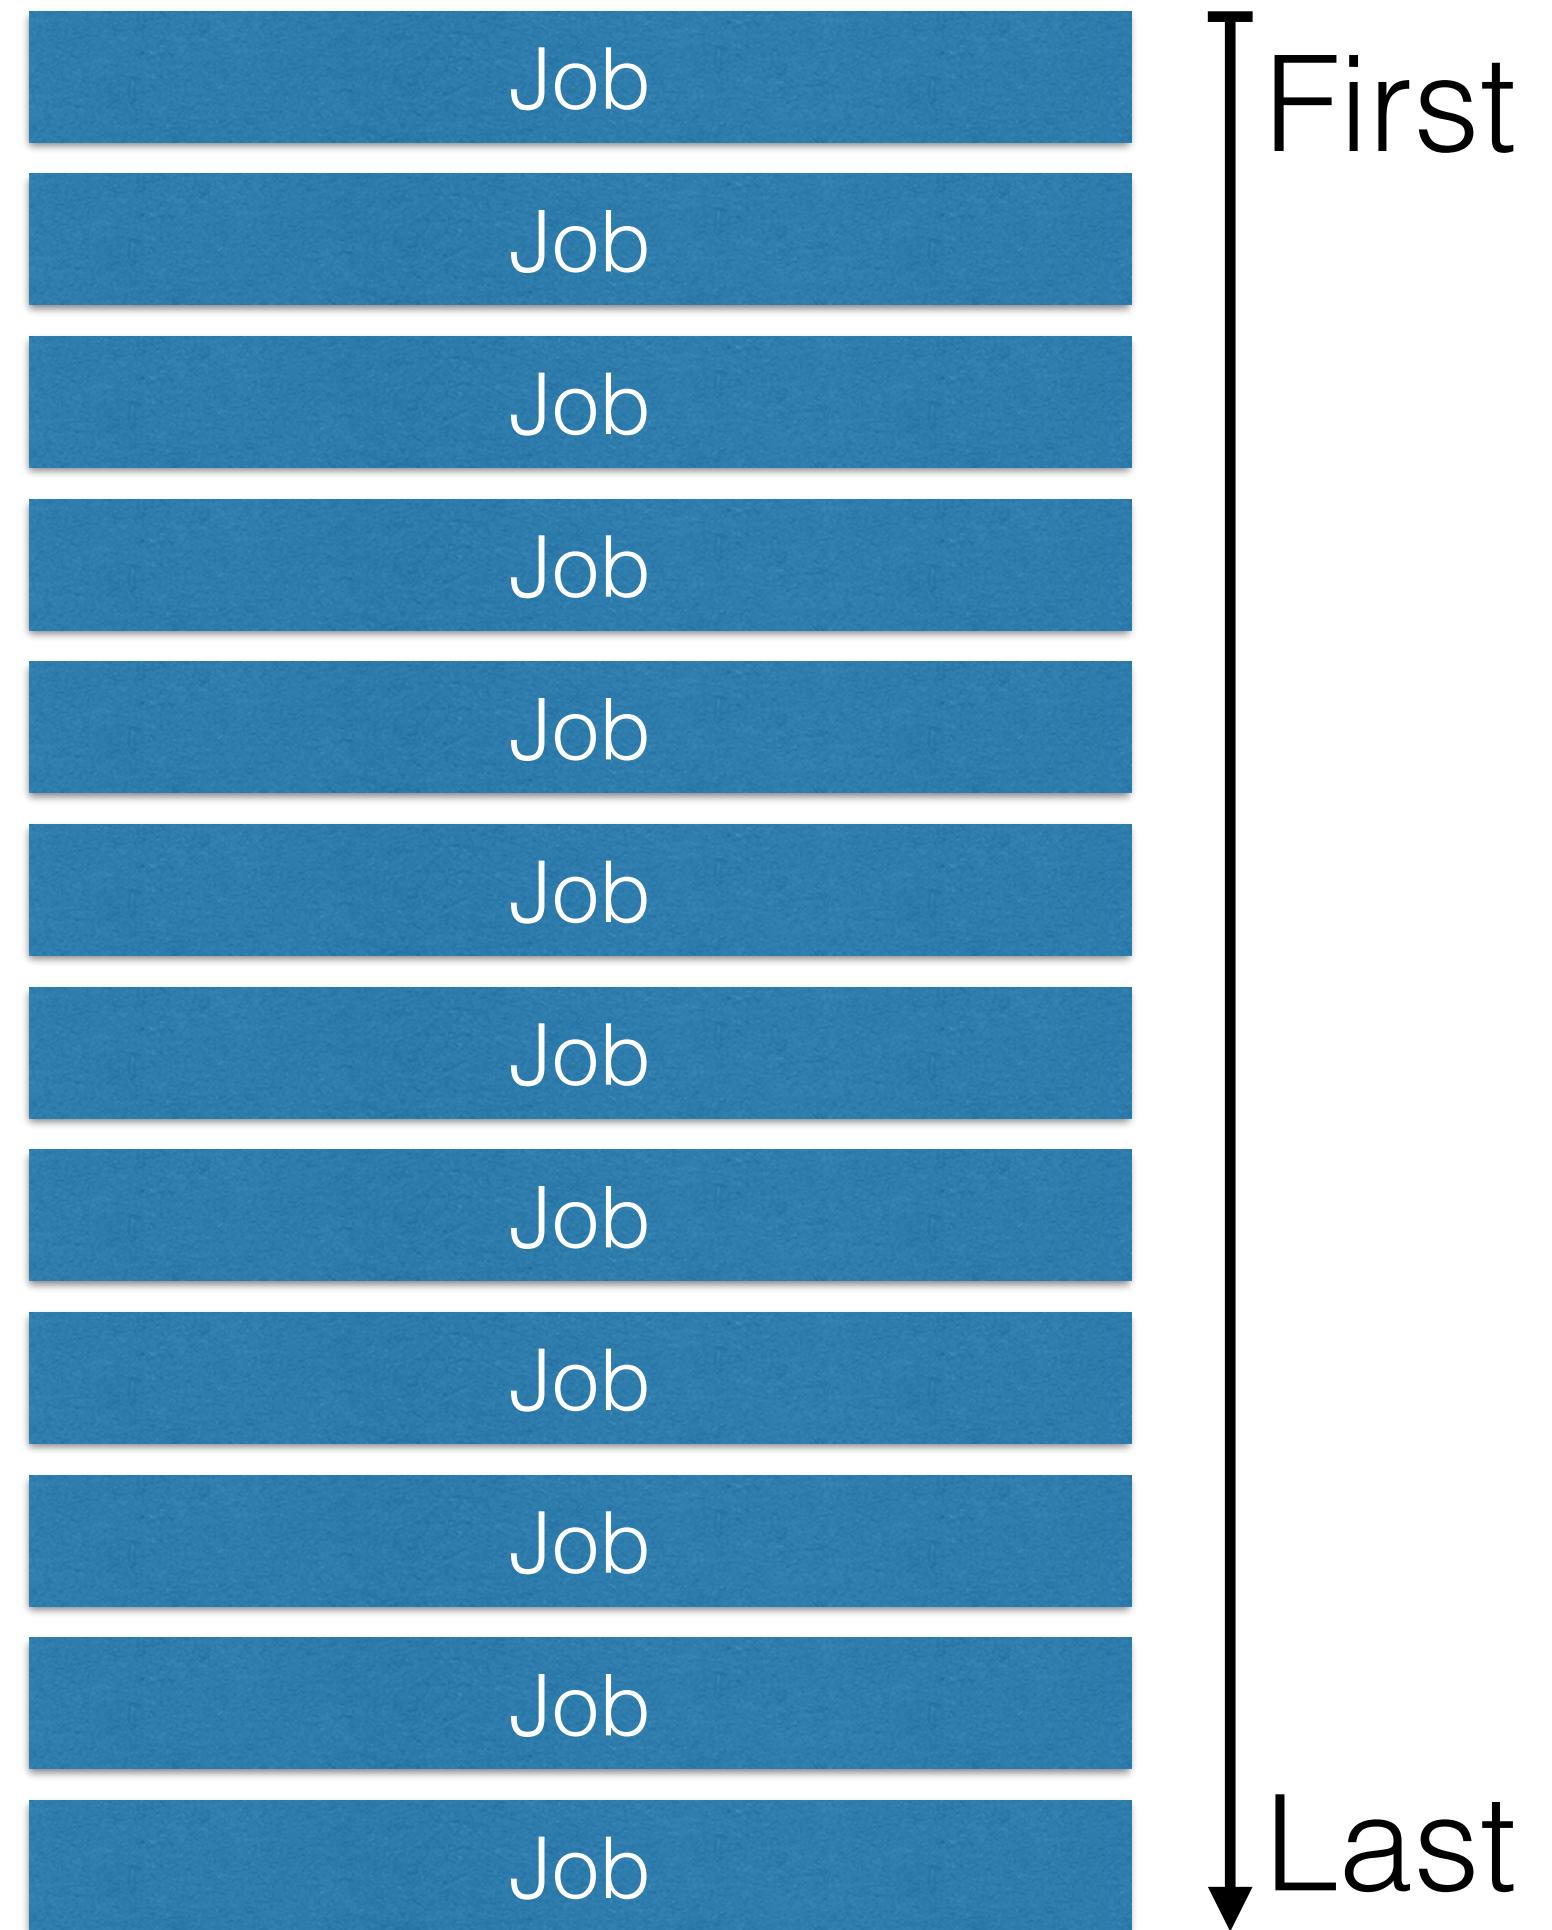
- Compute resources are shared by all users
- Queue systems manage access to specific resources
- There are many queue systems around: Moab, Torque, LoadLeveler, Condor, Oracle Grid Engine, Argent Job Scheduler, Platform LSF, etc.
- ManeFrame uses Simple Linux Utility for Resource Management (SLURM)





Resource Management

- ManeFrame resource allocation is governed by a fair-share factor
 - Preference is given to those who have submitted the fewest jobs
 - Prevents a single user from “blocking” subsequent job submissions of other users
- Allocations can be requested interactively or as a batch job
- Queues (or partitions) are used to distinguish resources by type and the maximum run time





ManeFrame II Partitions (Queues)

Queue	Quantity [Nodes]	Exclusivity	Duration
development	2	Partial	2 hours
htc	52	Shared	1 day
standard-mem-s	80	Exclusive	1 day
standard-mem-m	24	Exclusive	1 week
standard-mem-l	18	Exclusive	1 month
medium-mem-1-s	20	Exclusive	1 day
medium-mem-1-m	10	Exclusive	1 week
medium-mem-1-l	5	Exclusive	1 month
medium-mem-2	4	Exclusive	2 weeks
high-mem-1	5	Exclusive	2 weeks
high-mem-2	6	Exclusive	2 weeks
mic	36	Exclusive	1 week
gpgpu-1	36	Exclusive	1 week
v100x8	3	Shared	1 week
dcv	5	Exclusive	1 day



Basic SLURM Commands

Command	Description	Example Usage
sinfo	Displays partition and node state information	<code>sinfo</code>
squeue	Displays queue state information	<code>squeue -u \$USER</code>
sbatch	Submits batch script	<code>sbatch job.submit</code>
srun	Requests resources for interactive use	<code>srun -p development --exclusive --pty \$SHELL</code>
scancel	Cancel jobs	<code>scancel 12345678</code>

The background image shows a server room with two rows of tall, dark grey server racks. The racks have a perforated front panel. Above the racks, a complex network of orange, blue, and yellow cables is suspended from a metal truss ceiling. The ceiling has several rectangular light fixtures. The floor is a light-colored tile.

Need help or have questions?

rkalescky@smu.edu