

Python Workflows on ManeFrame II (M2)

Robert Kalescky

HPC Applications Scientist

March 18, 2021

Research and Data Sciences Services

Office of Information Technology

Center for Research Computing

Southern Methodist University



Research Support

ManeFrame II (M2)

ManeFrame II (M2) HPC OnDemand Web Portal

Slurm

Python Environments

Containers

Research Support



- Provide research computing support, consultations, and collaborations
- Data Science - Dr. Eric Godat
- High-Performance Computing - Dr. Robert Kalescky
- Custom Devices (IOT, wearables, etc.) - Guillermo Vasquez



- Maintains our primary shared resource for research computing, ManeFrame II (M2), in collaboration with OIT
- Provides research computing tools, support, and training to all faculty, staff, and students using research computing resources
- www.smu.edu/crc has documentation and news
- help@smu.edu or rkalescky@smu.edu for help
- Request an account at www.smu.edu/crc



Date	Workshop
March 4	ManeFrame II (M2) Introduction
March 11	Intel oneAPI HPC Toolkit Overview
March 18	Python Workflows on ManeFrame II (M2)
March 25	NVIDIA HPC SDK Overview
April 1	ManeFrame II (M2) Introduction
April 8	R Workflows on ManeFrame II (M2)
April 15	Introduction to OpenMP and OpenACC
April 22	Introduction to MPI
April 29	Introduction to Standard C++ Parallel Algorithms

Table 1: Workshops will be held each Thursday via Zoom from 12:00 to 2:00 PM. Register to receive Zoom coordinates. Sessions will be recorded and posted along with session materials. [Register!](#)

ManeFrame II (M2)



Type	Quantity	Cores	Memory [GB]	Additional Resources
Standard-Memory	176	36	256	
Medium-Memory-1	35	36	768	
Medium-Memory-2	4	24	768	3 TB SSD local scratch
High-Memory-1	5	36	1,536	
High-Memory-2	6	40	1,536	3 TB SSD local scratch
GPGPU-1	36	36	256	NVIDIA P100 GPU has 3,584 CUDA cores and 16 GB CoWoS
MIC-1	36	64	384	16 GB of high bandwidth (400 GB/s) stacked memory
VDI	5	36	256	NVIDIA Quadro M5000 GPU
v100x8	3	36	768	8 NVIDIA V100 GPUs with 5,120 CUDA cores and 32 GB CoWoS
Faculty Partner Nodes	3			Various research specific NVIDIA GPU configurations
ManeFrame II	354	11,276	120 TB	2.8 PB storage and InfiniBand network

ManeFrame II (M2) HPC OnDemand Web Portal



- Provides an integrated single access point for HPC resources on the ManeFrame II (M2) supercomputer
- Accessing the Portal:
 - Access to the HPC portal requires an existing M2 account
 - Go to **hpc.smu.edu**
 - Sign in using your SMU ID and SMU password
- HPC Portal Documentation
- HPC Portal Walkthrough Video



Slurm



```
1 import random, sys
2
3 def monte_carlo_pi(points):
4     return 4 * sum(1 for _ in range(points) if random.random()*2 +
5         ↪ random.random()*2 < 1) / float(points)
6
7 if __name__ == "__main__":
8     print(monte_carlo_pi(int(sys.argv[1])))
```

Listing 1: Serial algorithm to estimate the value of Pi.



```
1  import random, sys
2  import multiprocessing as mp
3
4  def check_point(points):
5      return sum(1 for _ in range(points) if random.random()*2 + random.random()*2
6         ↪ < 1)
7
8  def parallel_monte_carlo_pi(points, cores):
9      points_per_core = int(points / cores)
10     n = [points_per_core] * cores
11     n[0] = points_per_core + (points - (points_per_core * cores))
12     pool = mp.Pool(processes = cores)
13     results = pool.map(check_point, n)
14     return 4 * sum(results) / float(points)
15
16 if __name__ == "__main__":
17     print(parallel_monte_carlo_pi(int(sys.argv[1]), int(sys.argv[2])))
```

Listing 2: Parallel algorithm to estimate the value of Pi.



```
1 module load python
2 srun -p htc --mem=6G --pty $SHELL
```

Listing 3: Using `srun` to log into a compute node to run commands interactively.



```
1 srun -p htc --mem=6G python pi_monte_carlo.py 1000
```

Listing 4: Using `srun` to run commands directly on a compute node.



```
1 sbatch -p htc --mem=6G --wrap "sleep 30; time python pi_monte_carlo.py 1000"
```

Listing 5: Using `sbatch --wrap` wrap a commands in an `sbatch` script that is then submitted to the queue can run non-interactively.



```
1  #!/bin/bash
2  #SBATCH -J python
3  #SBATCH -o python_%j.out
4  #SBATCH -p htc
5  #SBATCH --mem=6G
6
7  module purge
8  module load python
9
10 time python pi_monte_carlo.py 1000
```

Listing 6: Using `sbatch` run serial computations via an `sbatch` script.



```
1  #!/bin/bash
2  #SBATCH -J pi
3  #SBATCH -o pi_%j.out
4  #SBATCH -p development
5  #SBATCH -N 1
6  #SBATCH --cpus-per-task=2
7  #SBATCH --mem=6G
8
9  module purge
10 module load python
11
12 time python pi_monte_carlo_shared.py 10000000 ${SLURM_CPUS_PER_TASK}
```

Listing 7: Using `sbatch` run parallel computations via an `sbatch` script.



```
1  #!/bin/bash
2  #SBATCH -J pi_array
3  #SBATCH -o pi_array_%a-%A.out
4  #SBATCH --array=1-4%2
5  #SBATCH -p development
6  #SBATCH --mem=6G
7
8  module purge
9  module load python
10
11  time python pi_monte_carlo.py $((100*${SLURM_ARRAY_TASK_ID}))
```

Listing 8: Using `sbatch --array` run parallel jobs via a single `sbatch` script.

Python Environments



- Python environments allow users to use specific versions of Python with the packages of their choice
- The packages can be installed via **conda** or **pip**, depending on the type of environment being used



- Anaconda environments are managed using **conda**, which is available via Anaconda installations such as **python/2** and **python/3** on ManeFrame II.



```
1 module purge
2 module load python/3
3 conda create -y -n jupyter_37 -c conda-forge jupyterlab python=3.7
4 source activate ~/.conda/envs/jupyter_37
5 which python3
6 python3 --version
7 deactivate
```

Listing 9: A specific version of Python is installed along with the JupyterLab package and its dependencies. The new environment is then loaded and then unloaded.



- Python 3 has native support for managing environments, which can be used with any Python 3 installation on ManeFrame II.



```
1 module purge
2 module load spack gcc-9.2
3 source <(spack module tcl loads --dependencies python@3.7%gcc@9.2)
4 python3 -m venv ~/.venv/jupyter_37
5 source ~/.venv/jupyter_37/bin/activate
6 pip3 install --upgrade pip
7 pip3 install --upgrade jupyterlab
8 which python3
9 python3 --version
10 deactivate
```

Listing 10: A specific installation of Python, in this case from Spack, is used and then JupyterLab and its dependencies are installed. The new environment is then loaded and then unloaded.

Containers



- Containers offer the ability to run fully customized software stacks, e.g. based on different Linux distributions and versions
- Containers are not virtual machines, where an entire hardware platform is virtualized, rather containers share a common kernel and access to physical hardware resources
- **Docker** is a popular container platform in development and server environments
- **Singularity** is a popular container platform in HPC environments
 - Docker is not directly support on ManeFrame II
 - Singularity can consume and run Docker containers



```
1  Bootstrap: docker                # Base container source
2  From: ubuntu:18.04              # Base container version
3
4  %post                            # Changes to container
5  export DEBIAN_FRONTEND=noninteractive
6  apt-get update
7  apt-get -y install\
8    python3-pip\
9    python3-numpy\
10   python3-pandas
11  pip3 install\
12   jupyterlab
13
14 %runscript                        # Default container command
15 python3
```

Listing 11: Example Singularity build file that uses Ubuntu 18.04 and Python3 with package installation via **apt** and **pip**.



```
1  #!/usr/bin/env sh
2
3  module purge
4  module load singularity/3.5.3
5
6  srun -p development -x k001 -c 1 --mem=1G\
7  --pty singularity build\
8  --fakeroot python3.sif python3.singularity
```

Listing 12: Steps to build a Singularity container image. Note that building Singularity container images from build scripts on M2 requires permission. Request permission via help@smu.edu.



```
1 FROM ubuntu:18.04
2
3 ENV DEBIAN_FRONTEND noninteractive
4
5 RUN apt-get update &&\
6     apt-get -y install\
7     python3-pip\
8     python3-numpy\
9     python3-pandas
10
11 RUN pip3 install\
12     jupyterlab
13
14 ENTRYPOINT ["python3"]
```

Listing 13: Example Dockerfile that uses Ubuntu 18.04 and Python3 with package installation via **apt** and **pip**.



```
1  #!/usr/bin/env sh
2
3  docker build -t python3:18.04 -f
   ↪ python3.dockerfile .
4  docker save -o python3_18_04.tar
   ↪ python3:18.04
5  scp python3_18_04.tar
   ↪ m2.smu.edu:~/workshops/examples/
```

Listing 14: Building the Docker container off of M2, exporting the container image, and uploading to M2.

```
1  #!/usr/bin/env sh
2
3  module purge
4  module load singularity/3.5.3.lua
5
6  srun -p development -x k001 -c 1
   ↪ --mem=1G --pty singularity build
   ↪ python3_18_04.sif docker-
   ↪ archive://python3_18_04.tar
```

Listing 15: Converting the uploaded Docker container image to a Singularity container image.



Need help or have questions?
rkalescky@smu.edu